

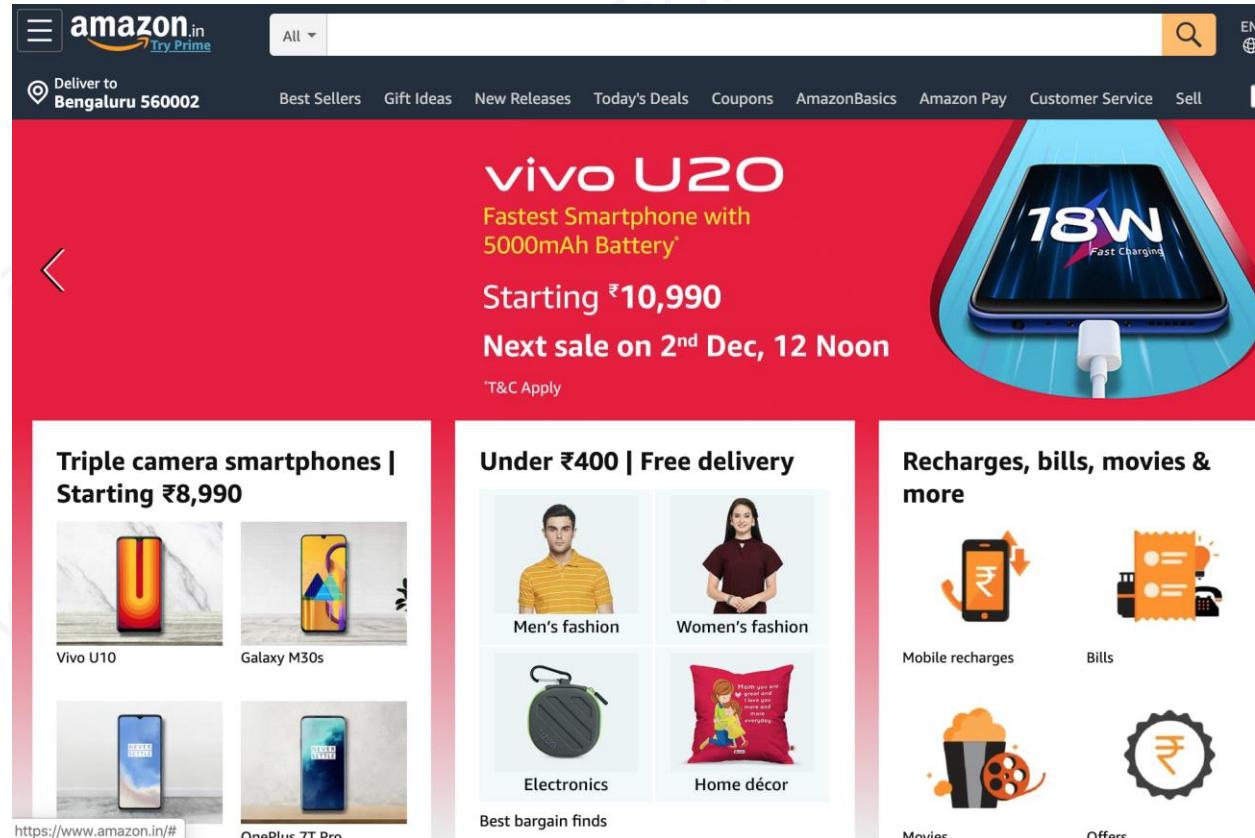


Developer To Architect

Moving from Application Design/Architecture to System Architecture

Our Objective

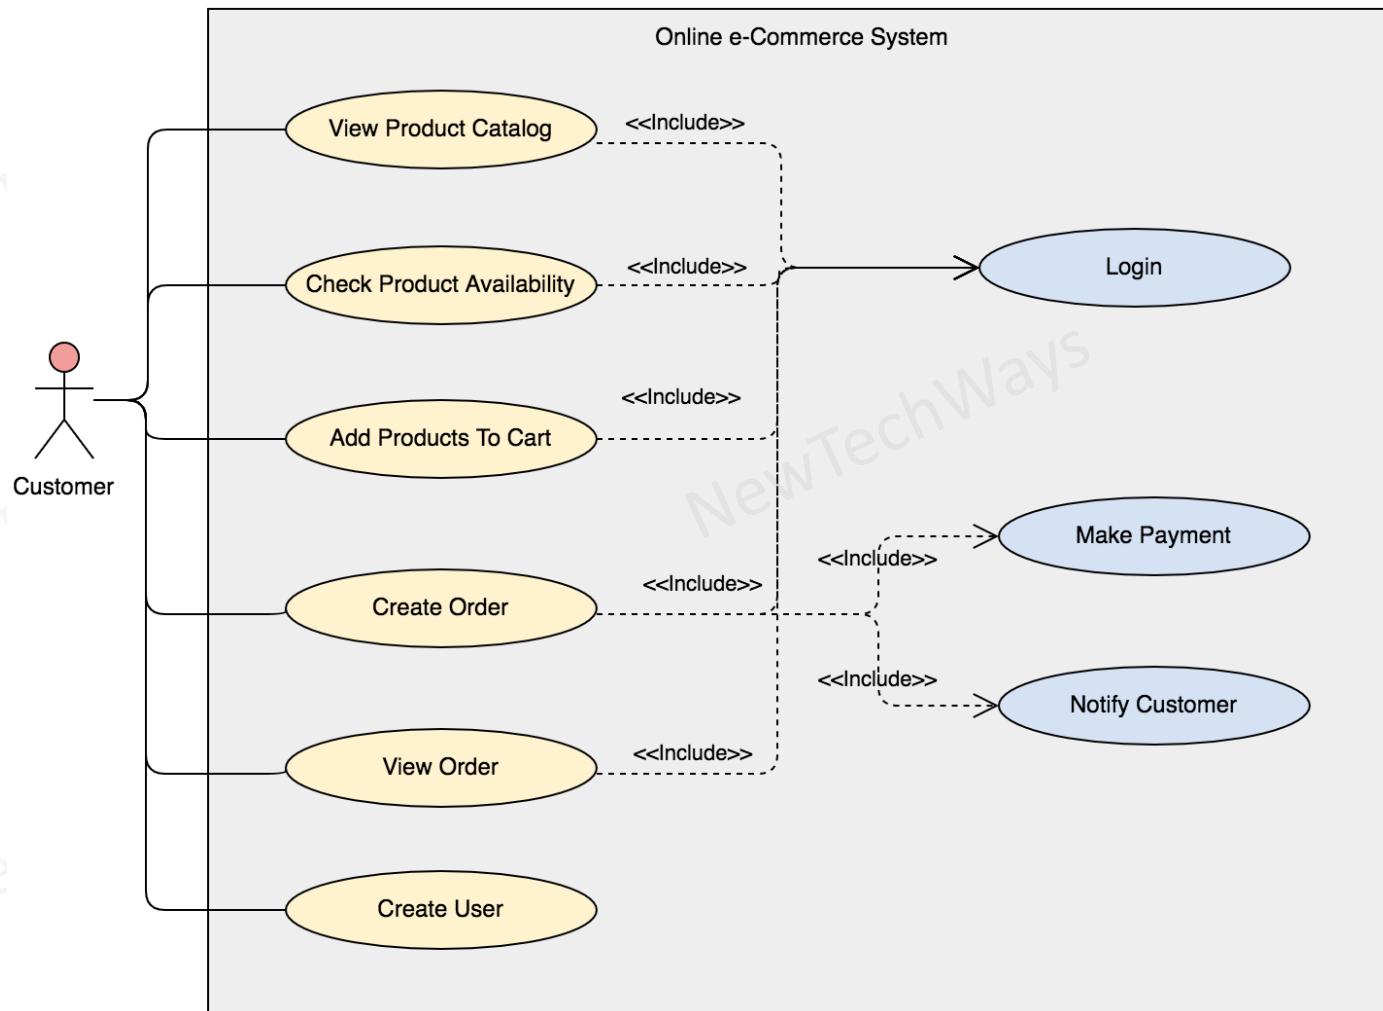
- Assume ‘Architecting’ a big e-commerce system as our objective



Application Architecture

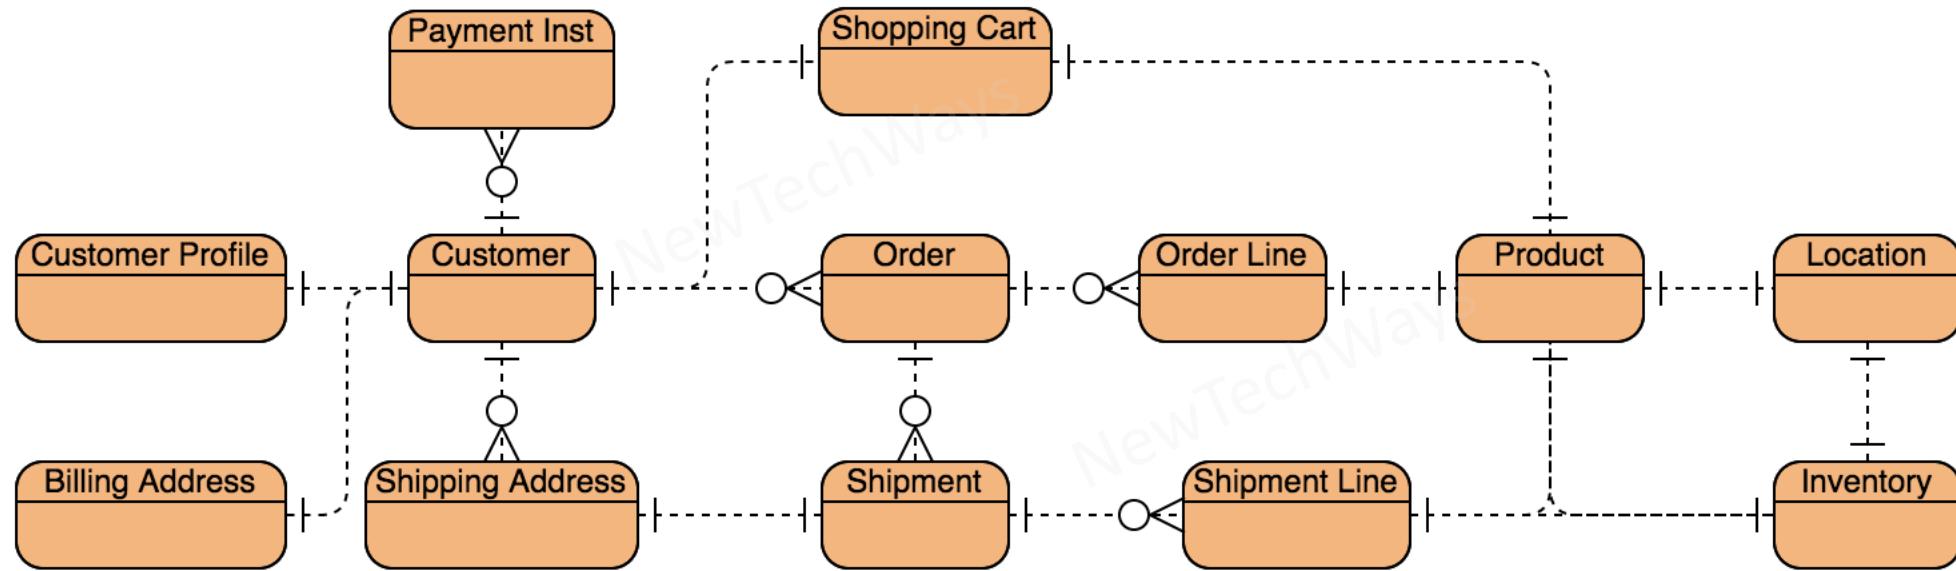
Use Case Model

4

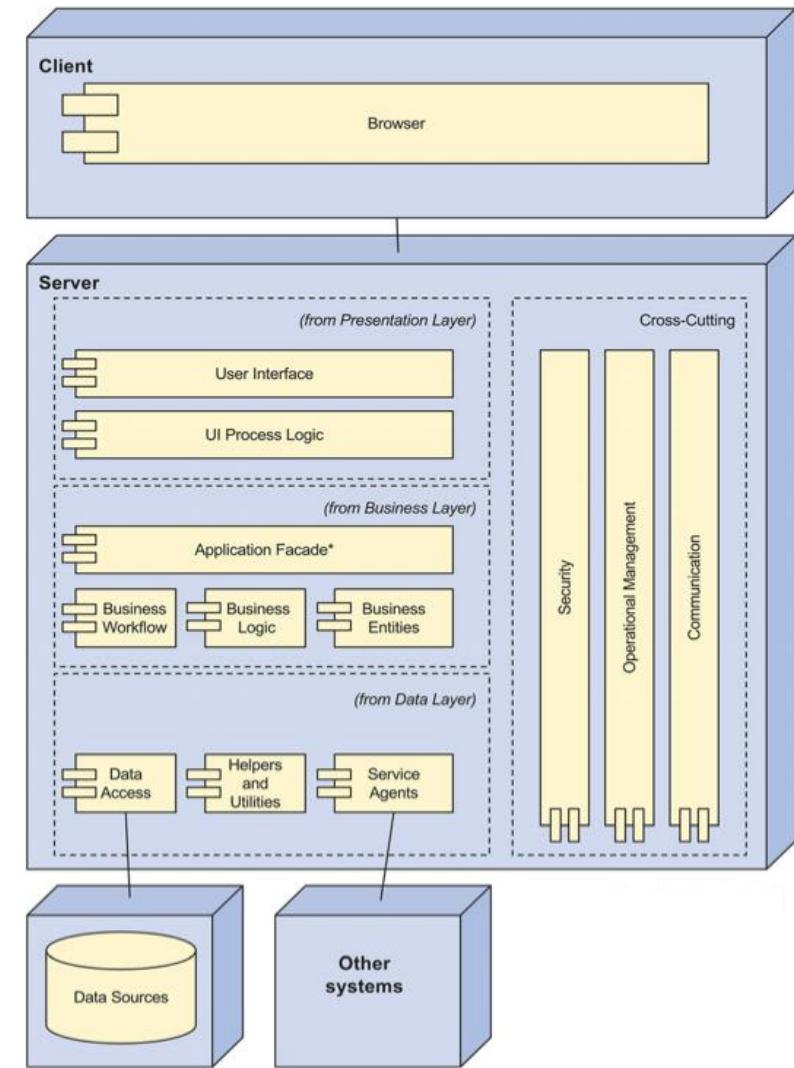
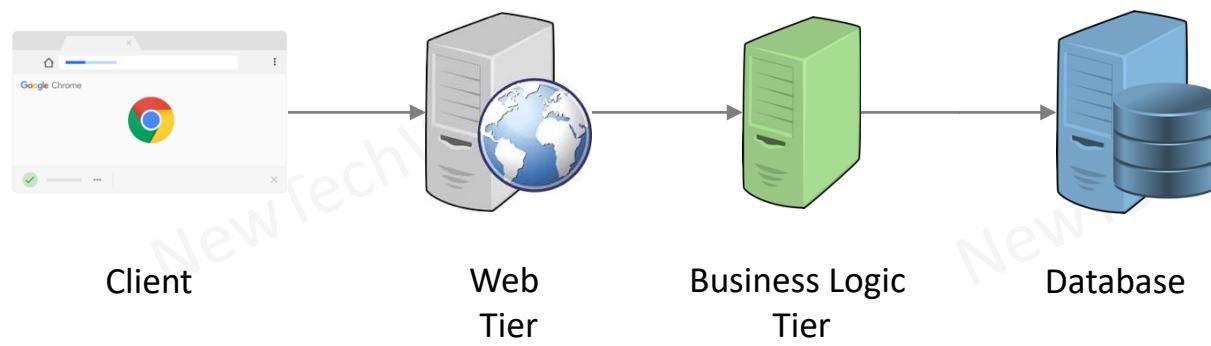


Domain Model

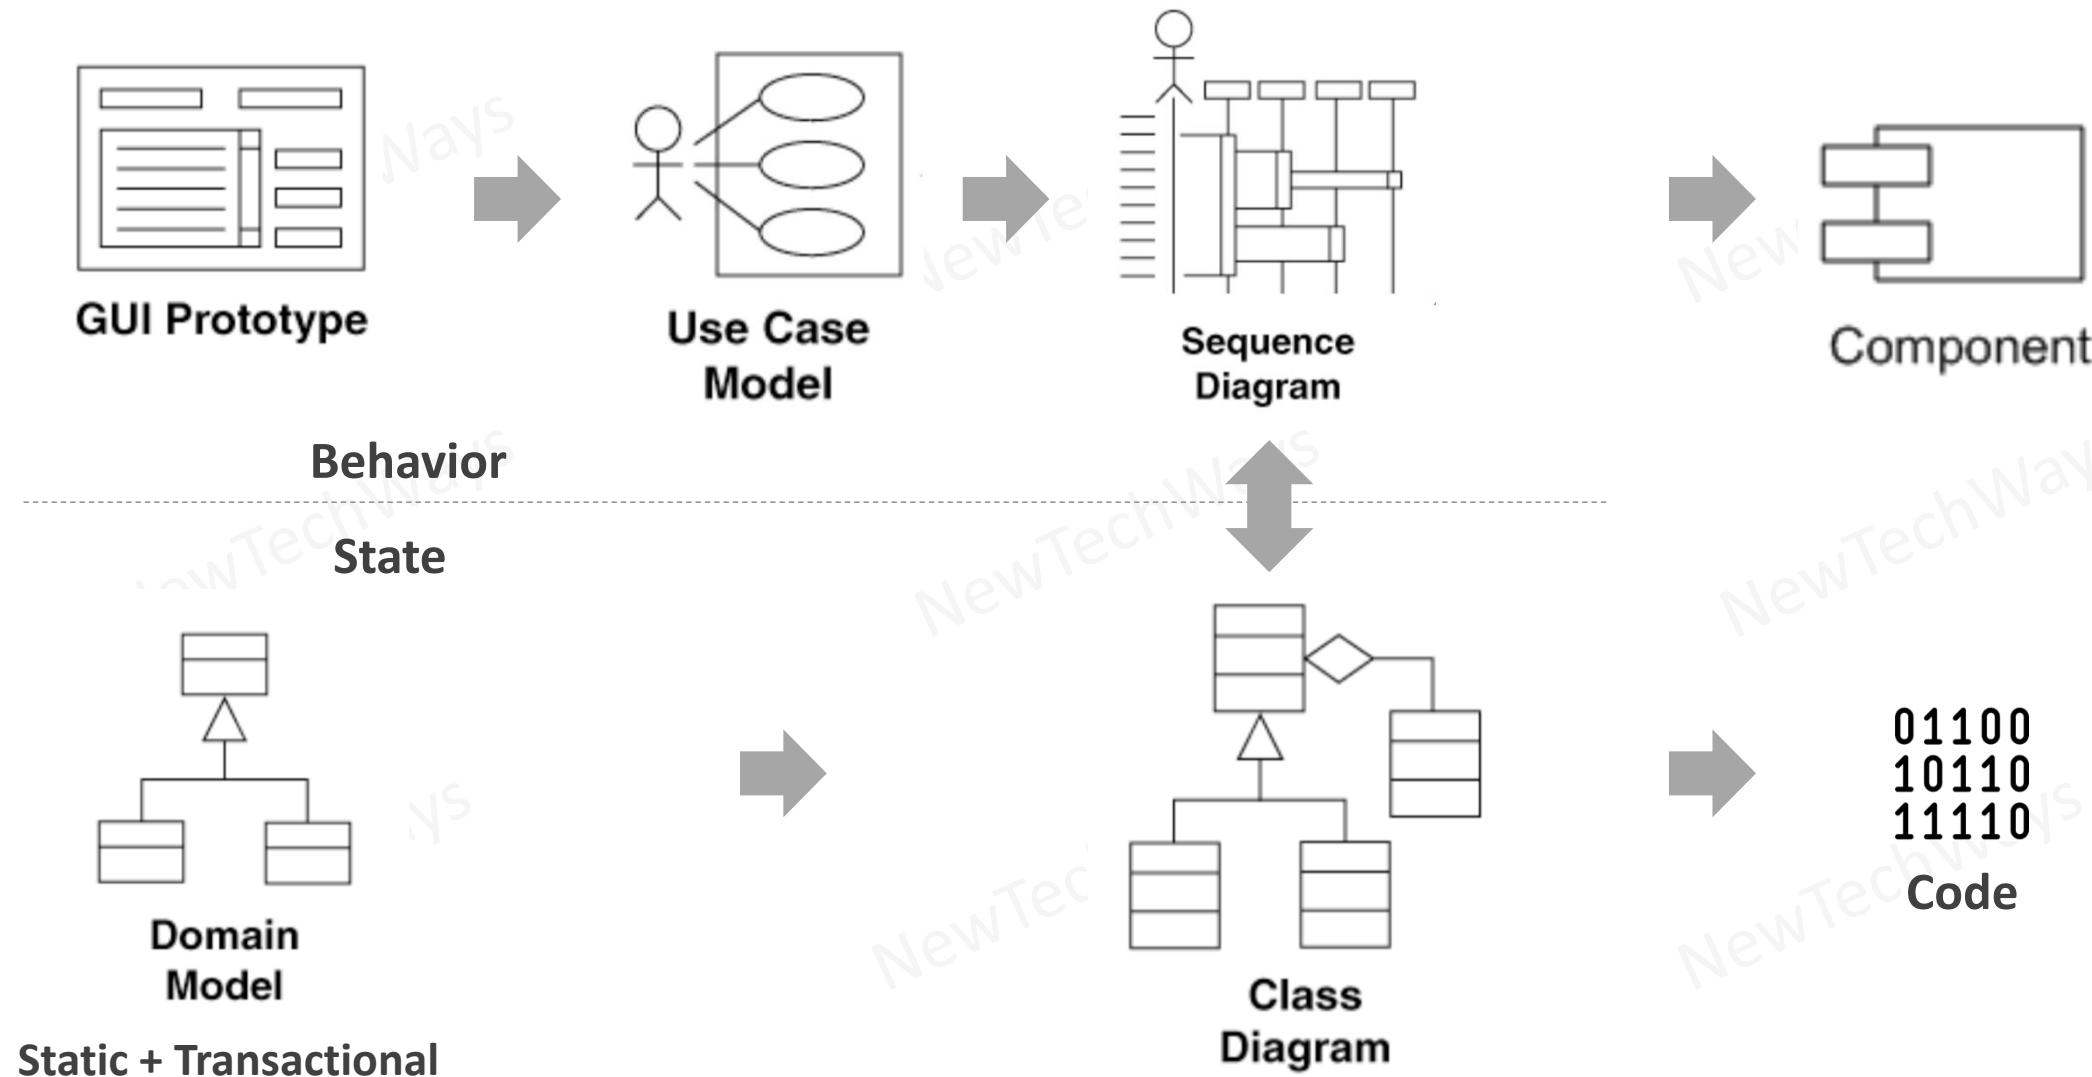
5



Component Model



Low Level Design



What About?

8



High Availability



Response Latency



Global Customers



System & Data Security



File storage for Catalog Images



Mobile Support



Unstructured Data Storage & Analytics



Cloud Deployment

Going Beyond

9



Application Architecture To System Architecture

Application Architecture Vs System Architecture

10

System level challenges
surface up in large scale systems



Scale, Reliability, Security, Deployment are
biggest concerns for a large-scale system



Latency Requirements

- UI page load in < 250ms
 - across the globe
- Customer transactions < 3 sec
- Latency is not an option
 - 10 years ago, Amazon found that every 100ms of latency cost them 1% in sales
 - Now Akamai study shows that every 100-millisecond delay in website load time can hurt sales by 6%
 - Google found an extra .5 seconds in search page generation time dropped traffic by 20%
 - A stockbroker can lose millions if their trading platform is 5ms behind its competition



As page load time goes from:

1s to 3s the probability of bounce **increases 32%**

1s to 5s the probability of bounce **increases 90%**

1s to 6s the probability of bounce **increases 106%**

1s to 10s the probability of bounce **increases 123%**

Scalability Requirements

- 10 million requests/day
- 1K to 100K simultaneous users
- Global customer base
- 100 million products
- Transaction data for last 5 years
- Petabytes of log data
- For 10 M requests with average response size of 10KB
 - $\text{Data Outflow} = 10M \times 10KB = 100 \text{ GB}$
- For 100 M orders/year with average order size of 10 KB, five years data
 - $\text{Data storage} = 5 \times 100M \times 10KB = 5 \text{ TB}$
- For 100 M products with average product description/image size of 1 MB
 - $\text{Data Storage} = 100M \times 1MB = 100 \text{ TB}$
- Logs generated and archived
 - $\text{Data Storage} \Rightarrow \text{in Petabytes}$
- Buying customers in 100 countries

Availability & Reliability Requirements

- Unavailability results in
 - Business loss
 - Reputation loss
- 99.95% Availability
 - Maximum cumulative disruption of 4 hours 22 minutes in a year
- 99.99999999% Durability for storage systems
 - Data once stored is practically never lost
- Disaster Recovery
 - Operations to continue even if a region goes down due to a natural calamity



Security Requirements

- Infrastructure protection
 - Network access
 - System access
 - Service access
- Data Protection
 - Data sensitivity classification
 - Protect data at rest
 - Protect data on wire
 - Data backup & replication
- Identity & Access Management
 - Authentication, Authorization
 - Role based access

“The knock-on effect of a data breach can be devastating for a company. When customers start taking their business—and their money—elsewhere, that can be a real body blow.”

CHRISTOPHER GRAHAM

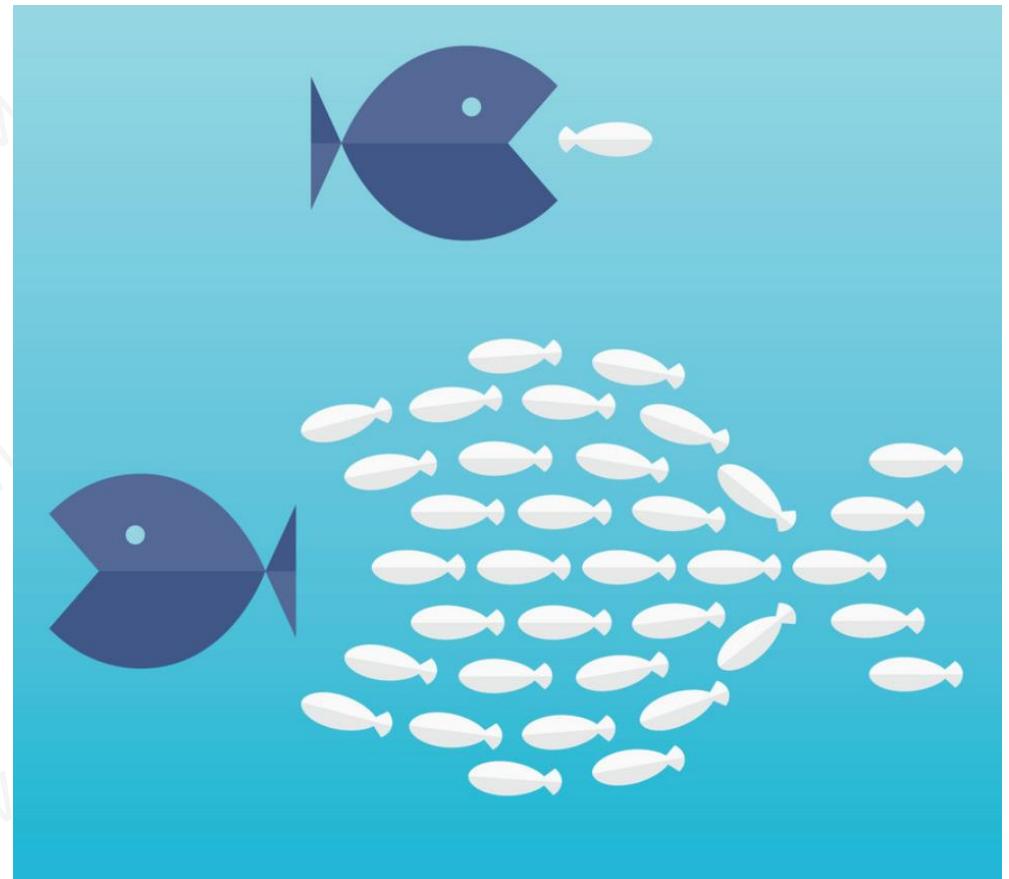


Designing System Architecture

Scalability Principle

16

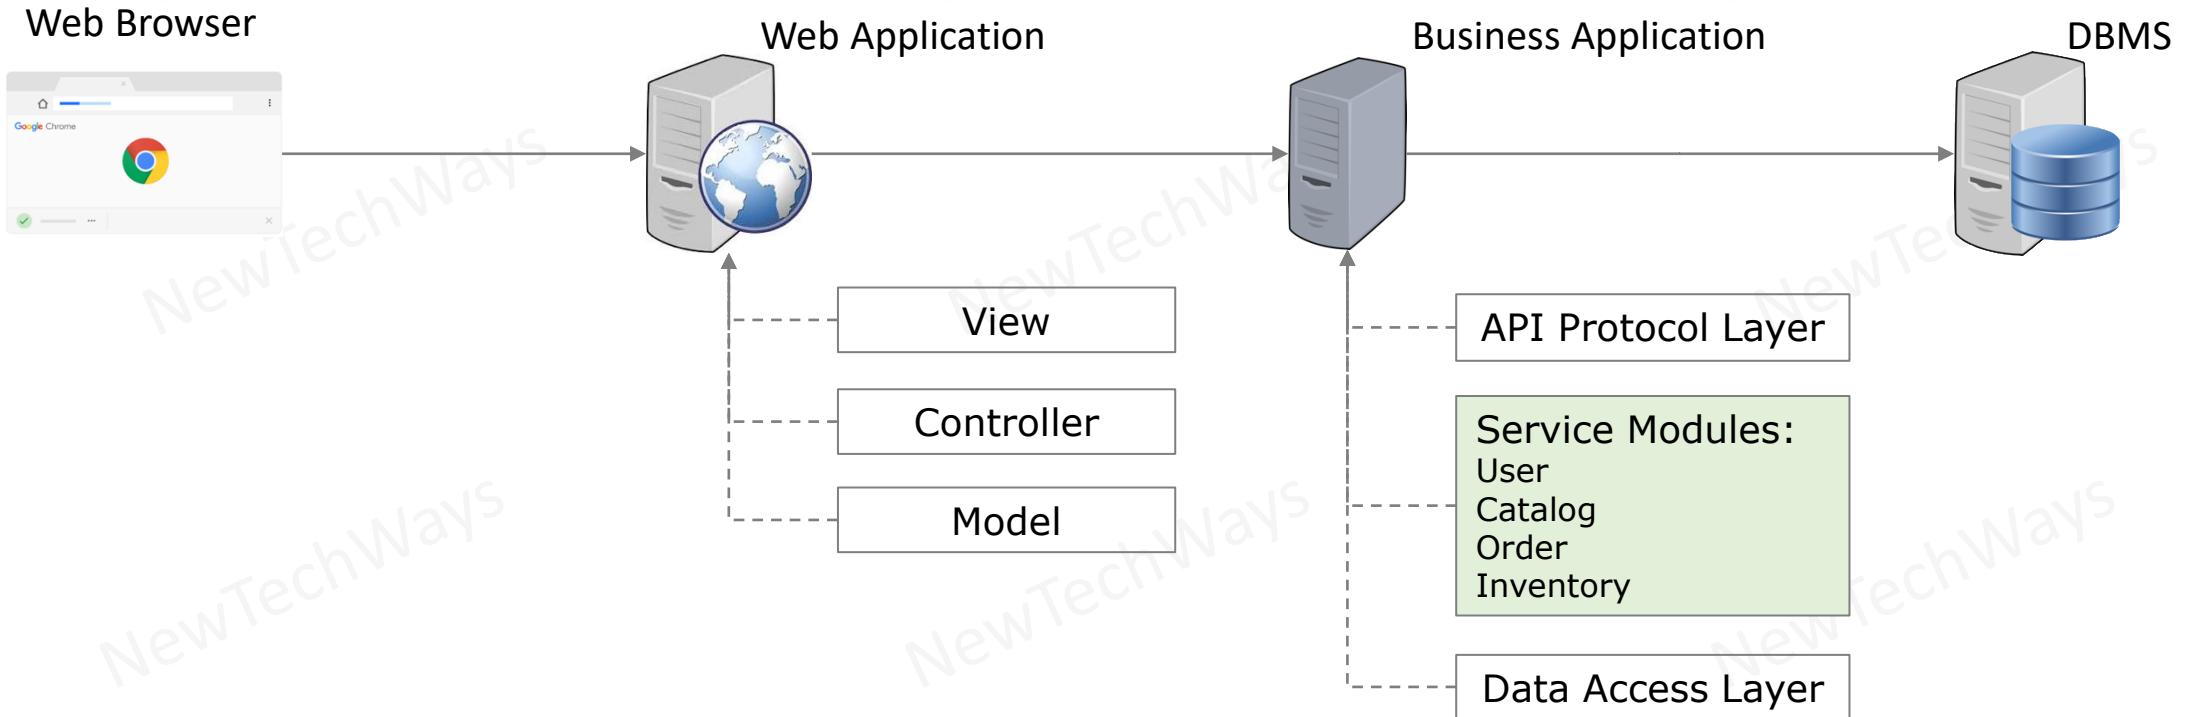
- Monolith is an anti-pattern for Scalability
- Scalability goes up with
 - Decentralization
 - More specialized workers – Services
 - More workers – Instances, Processors, Threads
 - Independence
 - Multiple workers are as good as a single worker if they can't work independently
 - They must work concurrently to maximum extent
 - Independence is impeded by
 - Shared resources
 - Shared mutable data



Modularity

17

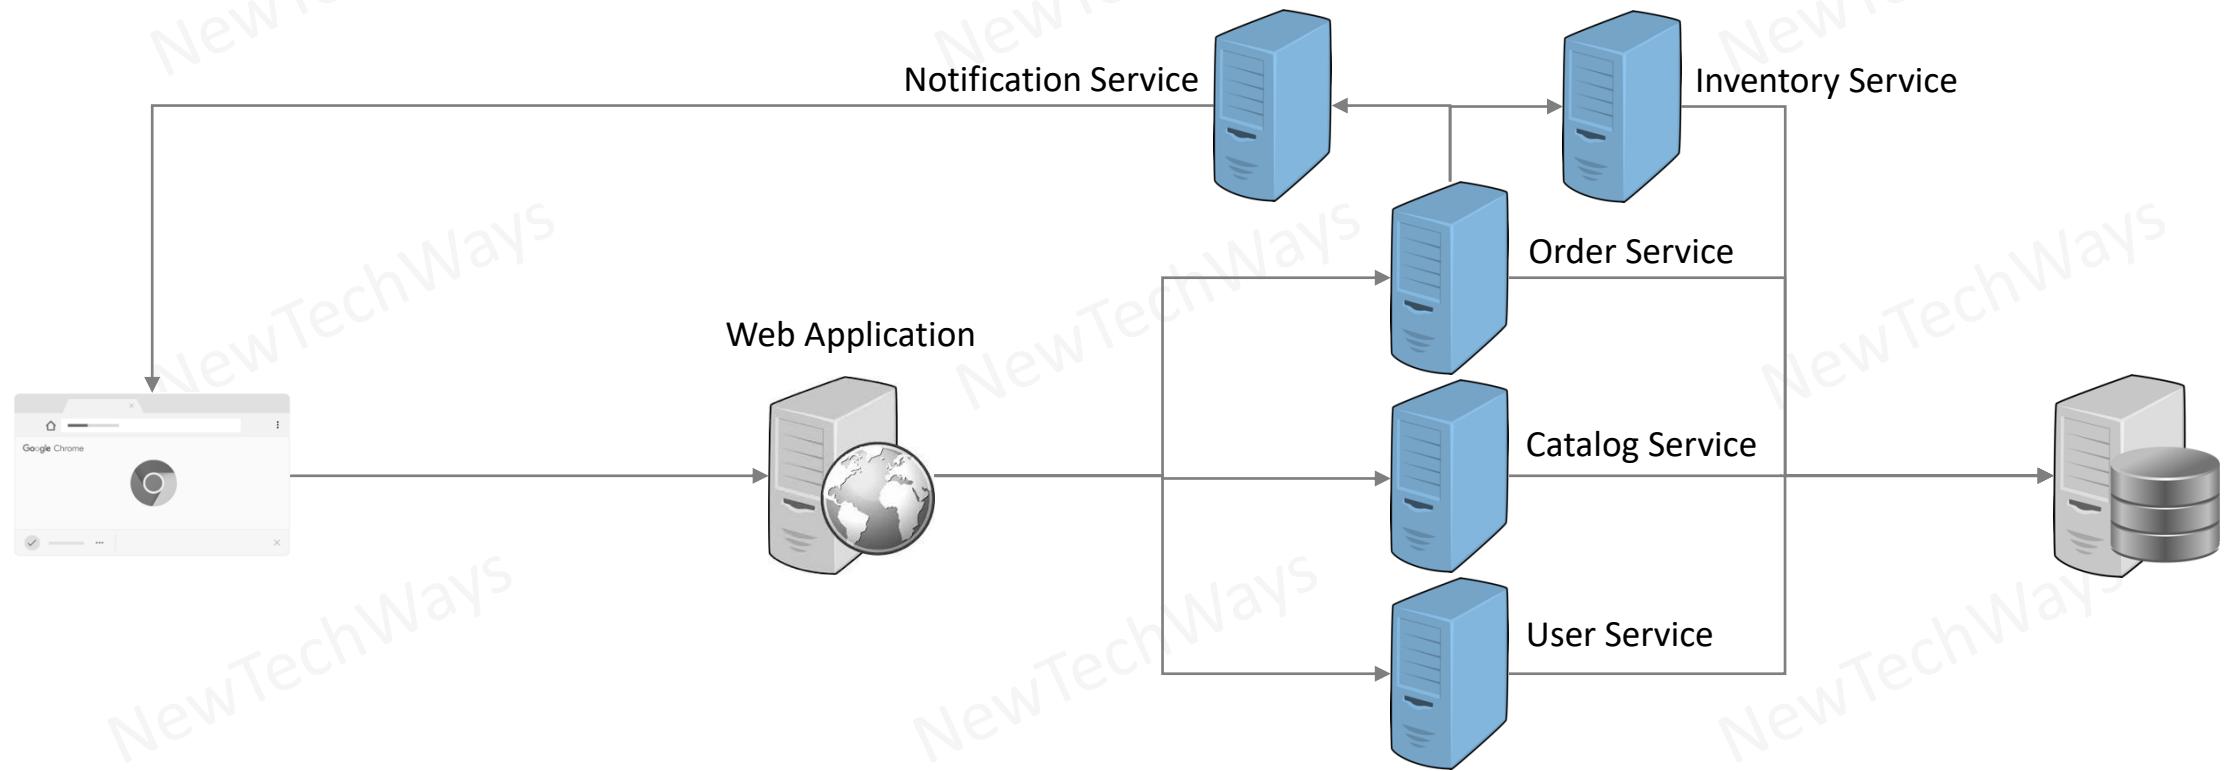
- Scalable architecture starts with modularity
 - Provides the foundation for breaking a system function/service into more specialized functions/services



Specialized Services – WebServices

18

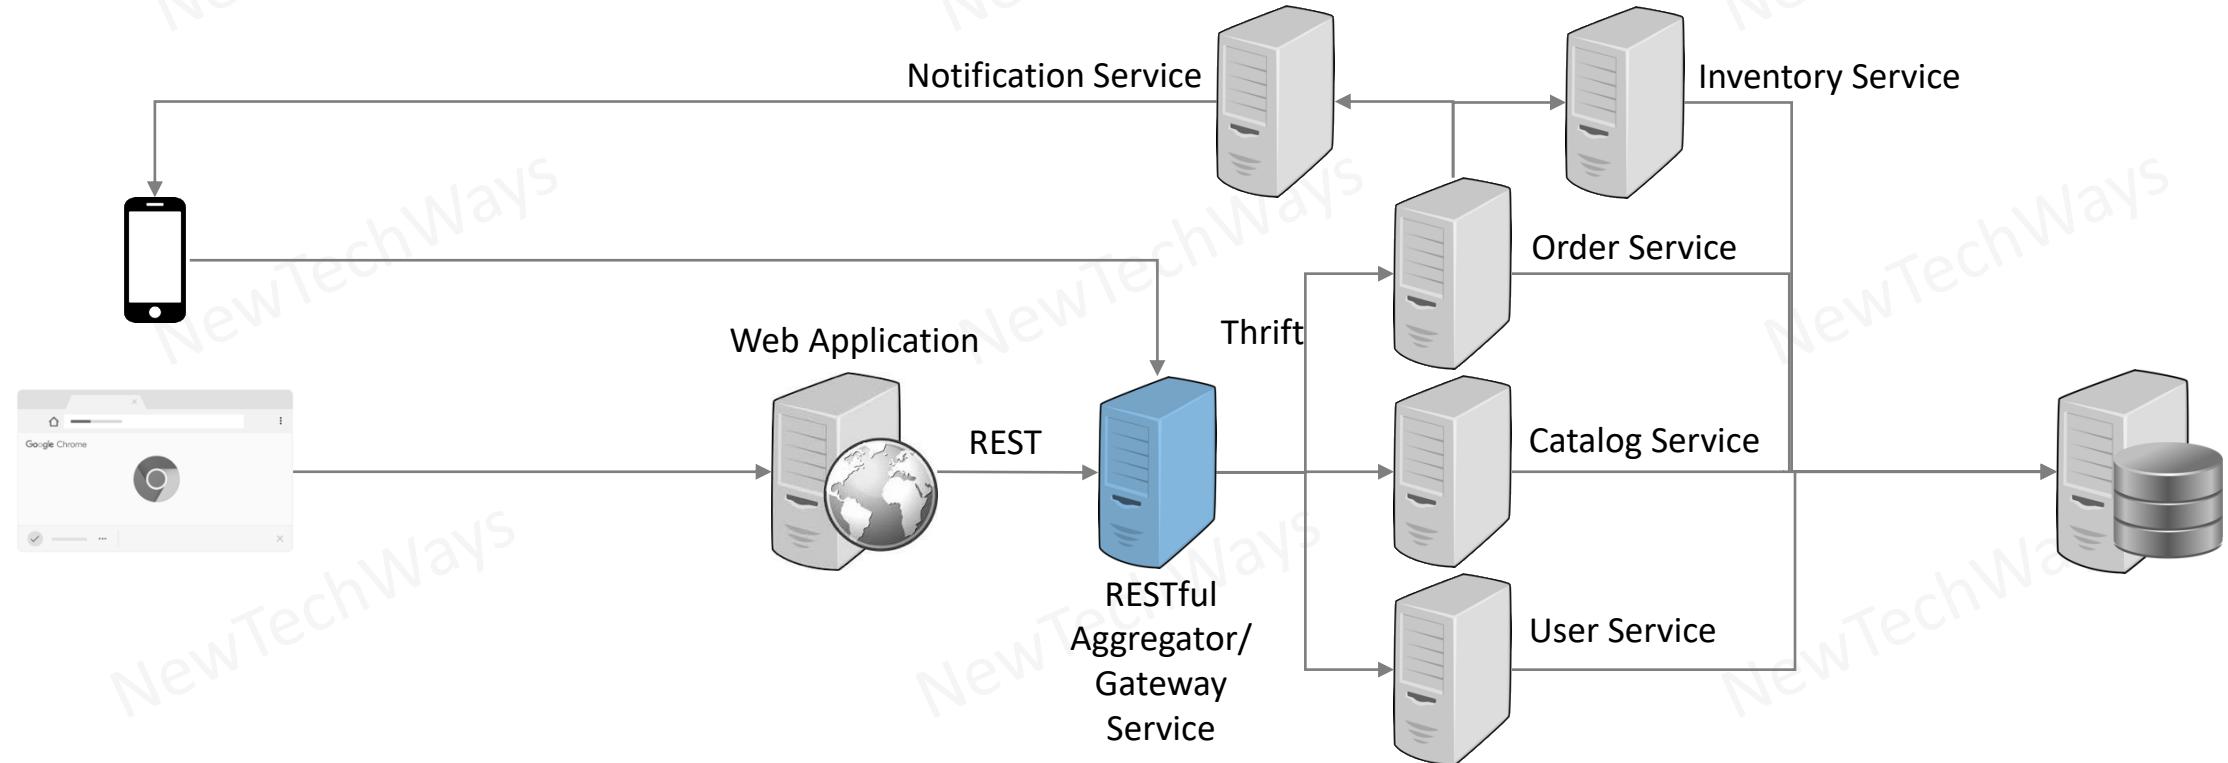
- Services can be scaled differently e.g., Number of instances



Aggregator Service & RESTful API – Mobile Support

19

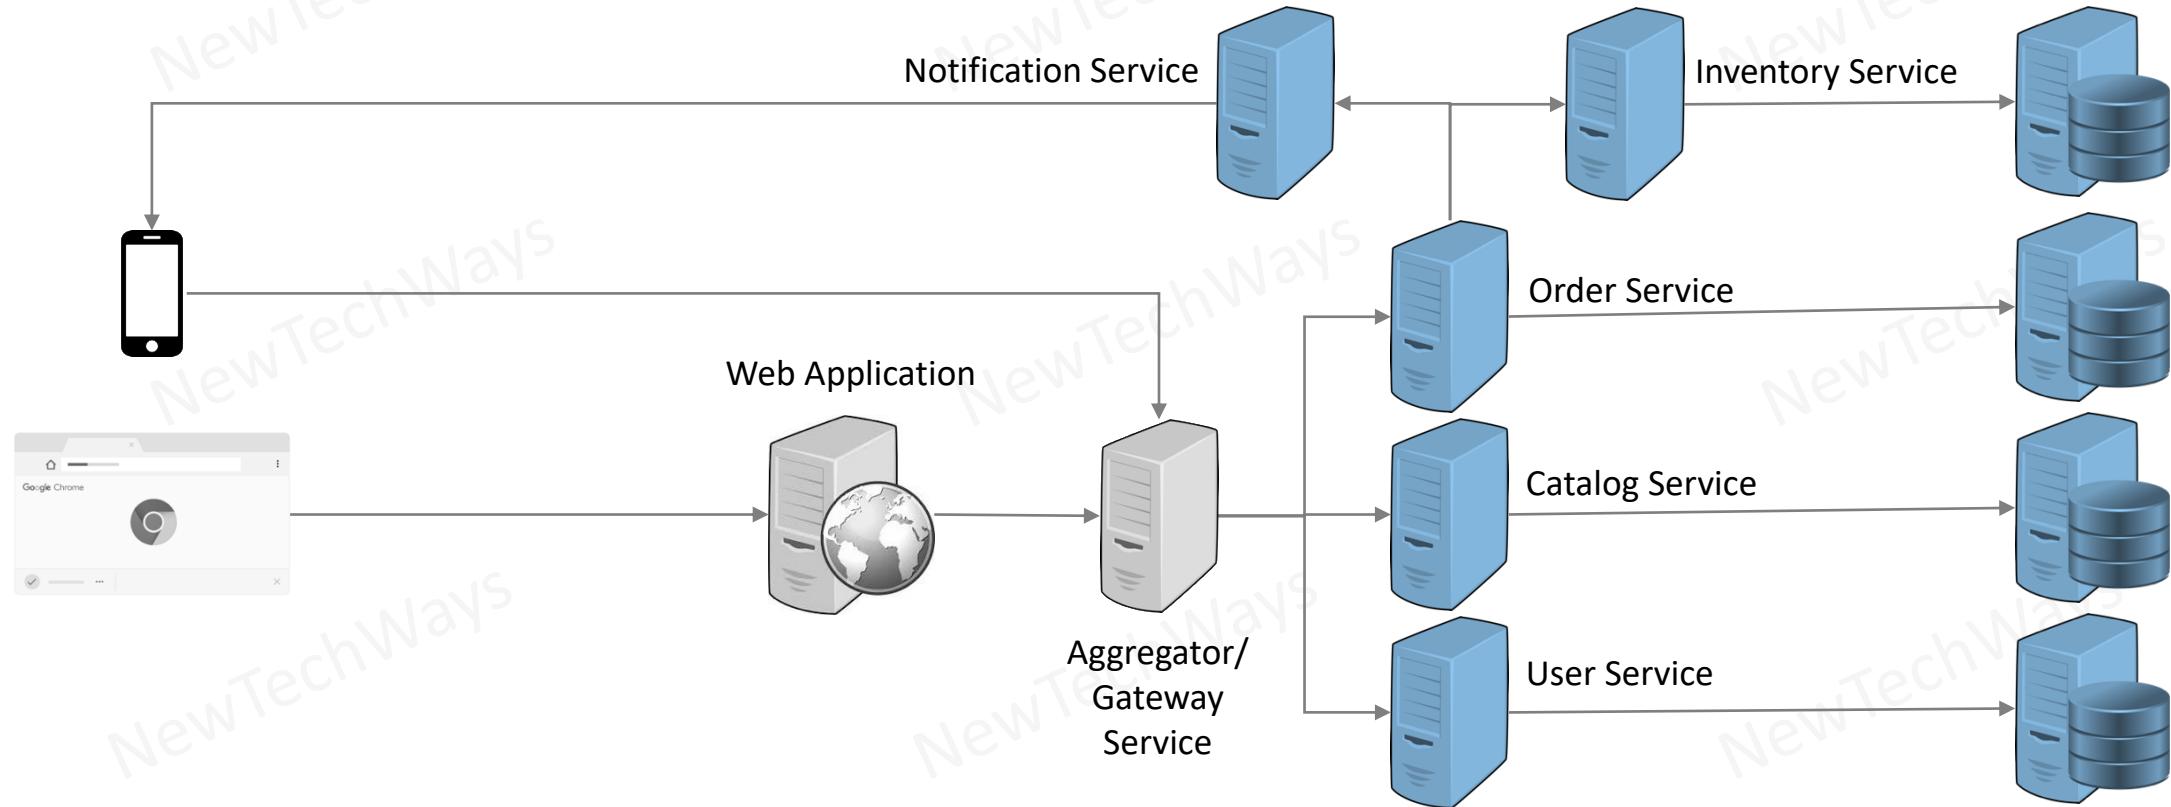
- REST for external interface interoperability & Mobile Support
- RPC for internal communication performance



Independent Services – MicroServices

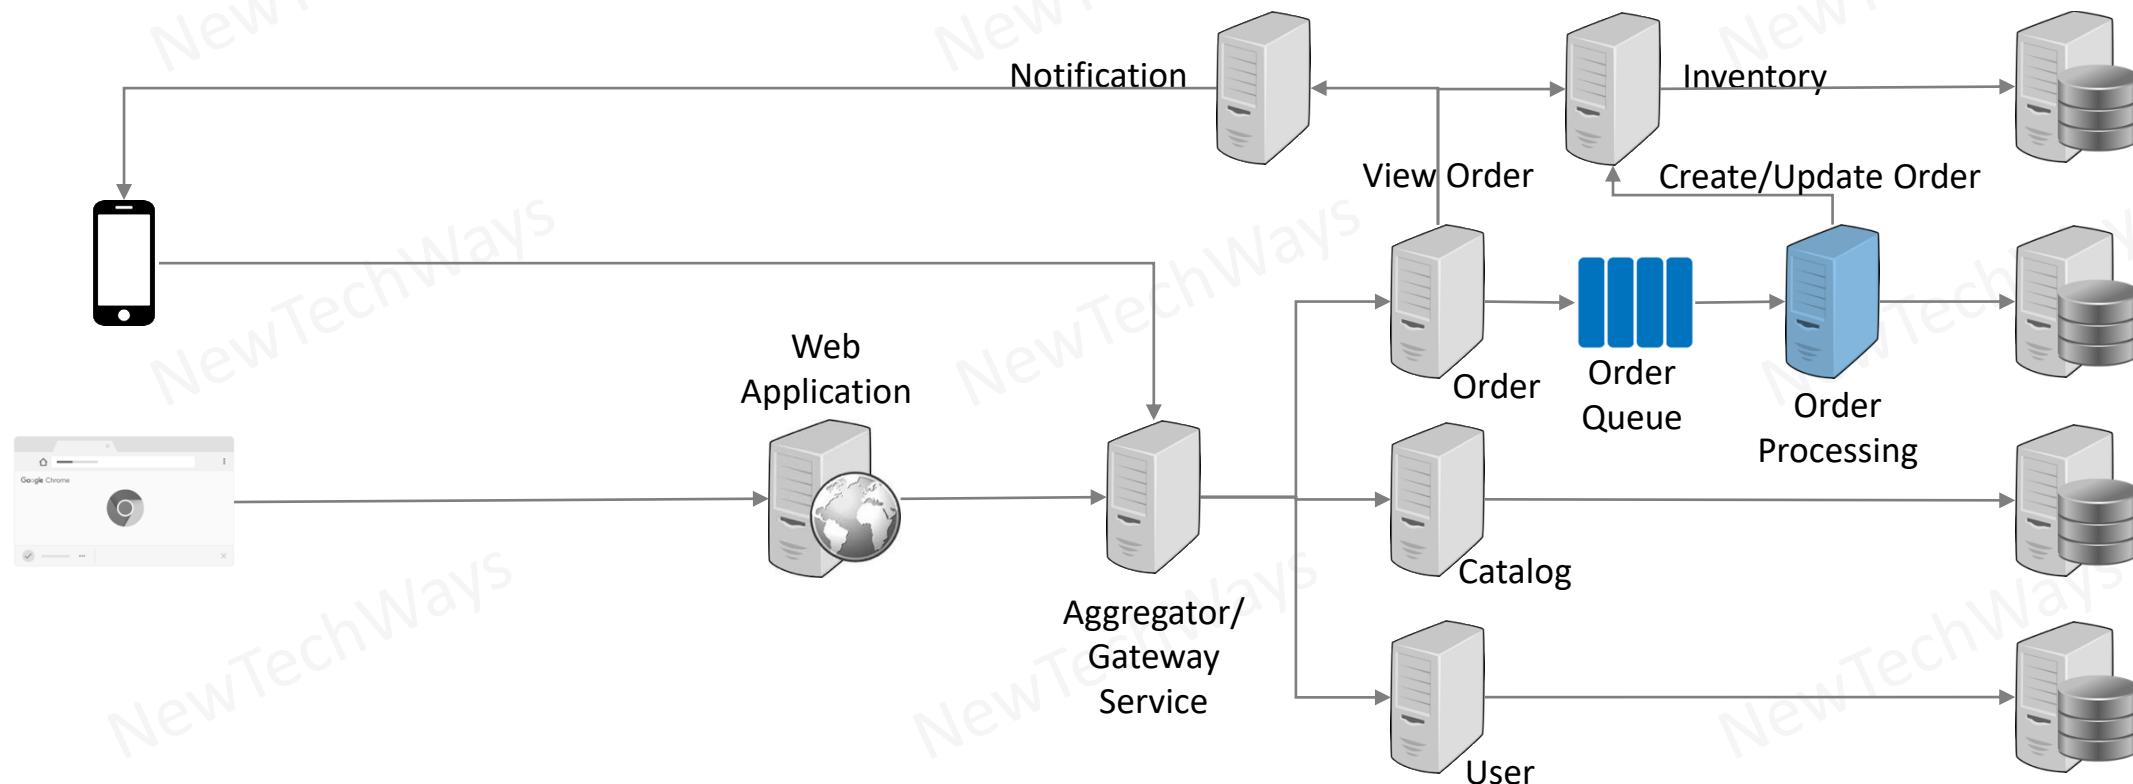
20

- Micro-Services can be scaled differently and deployed independently



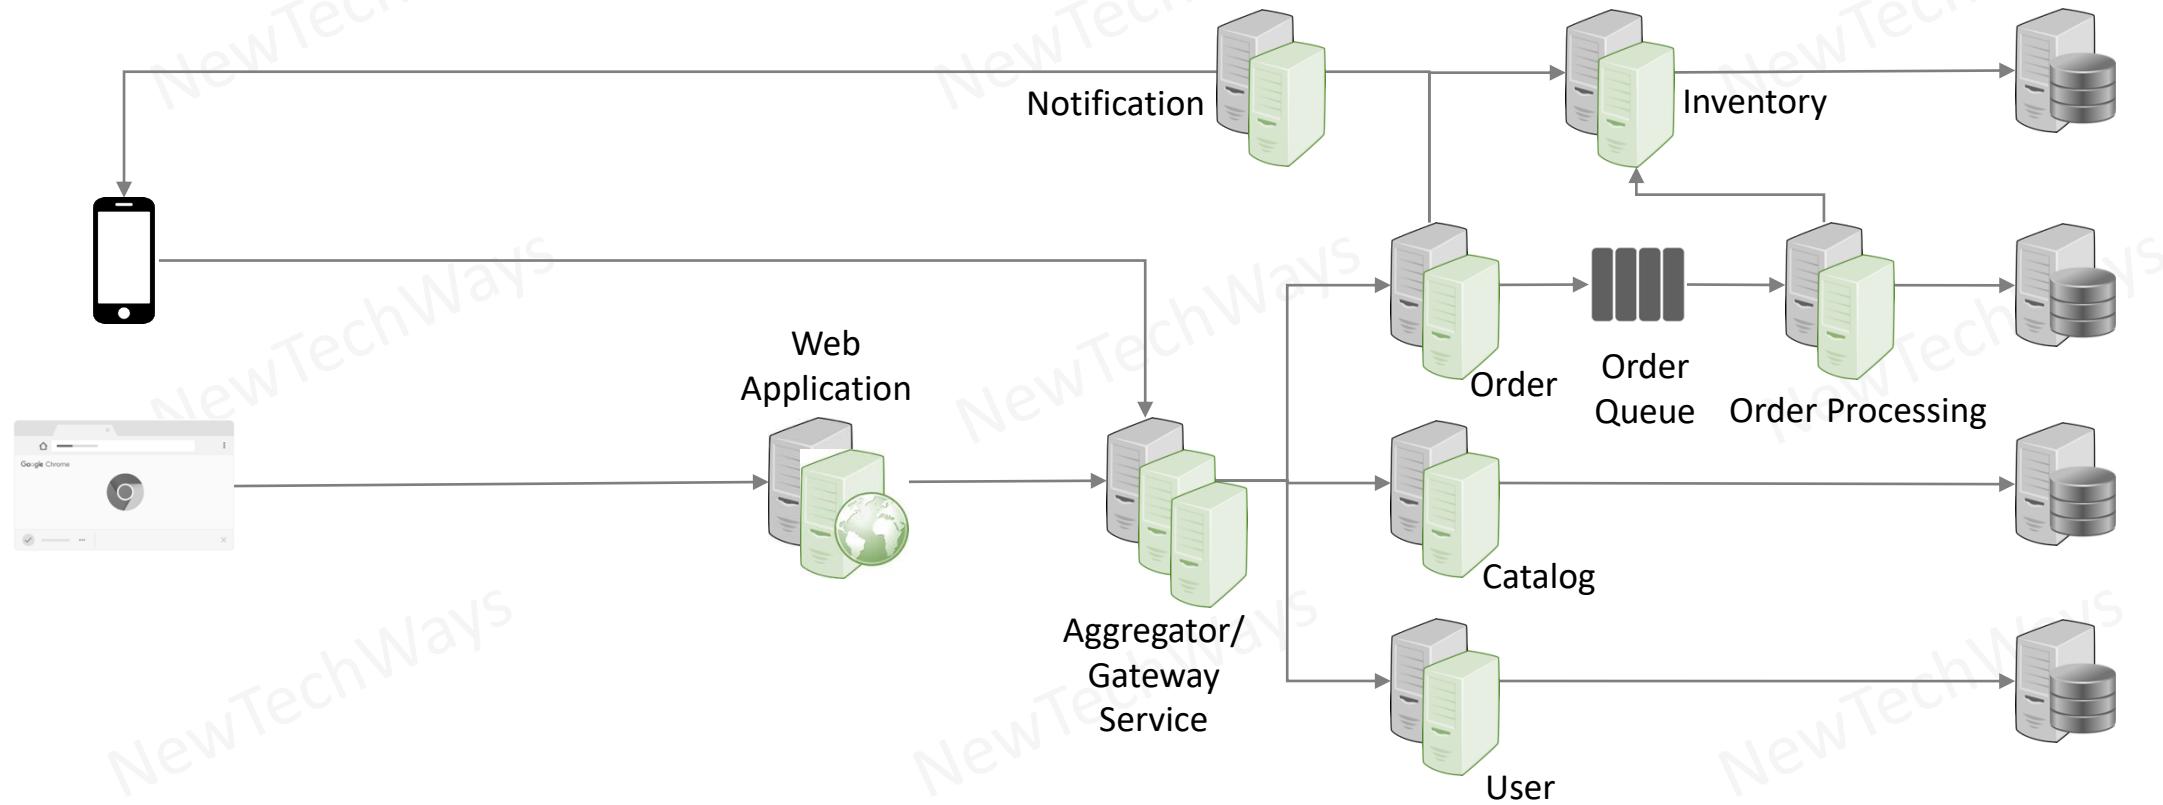
Asynchronous Services

- Updates can be done asynchronously to buffer peak loads and to reduce response latency



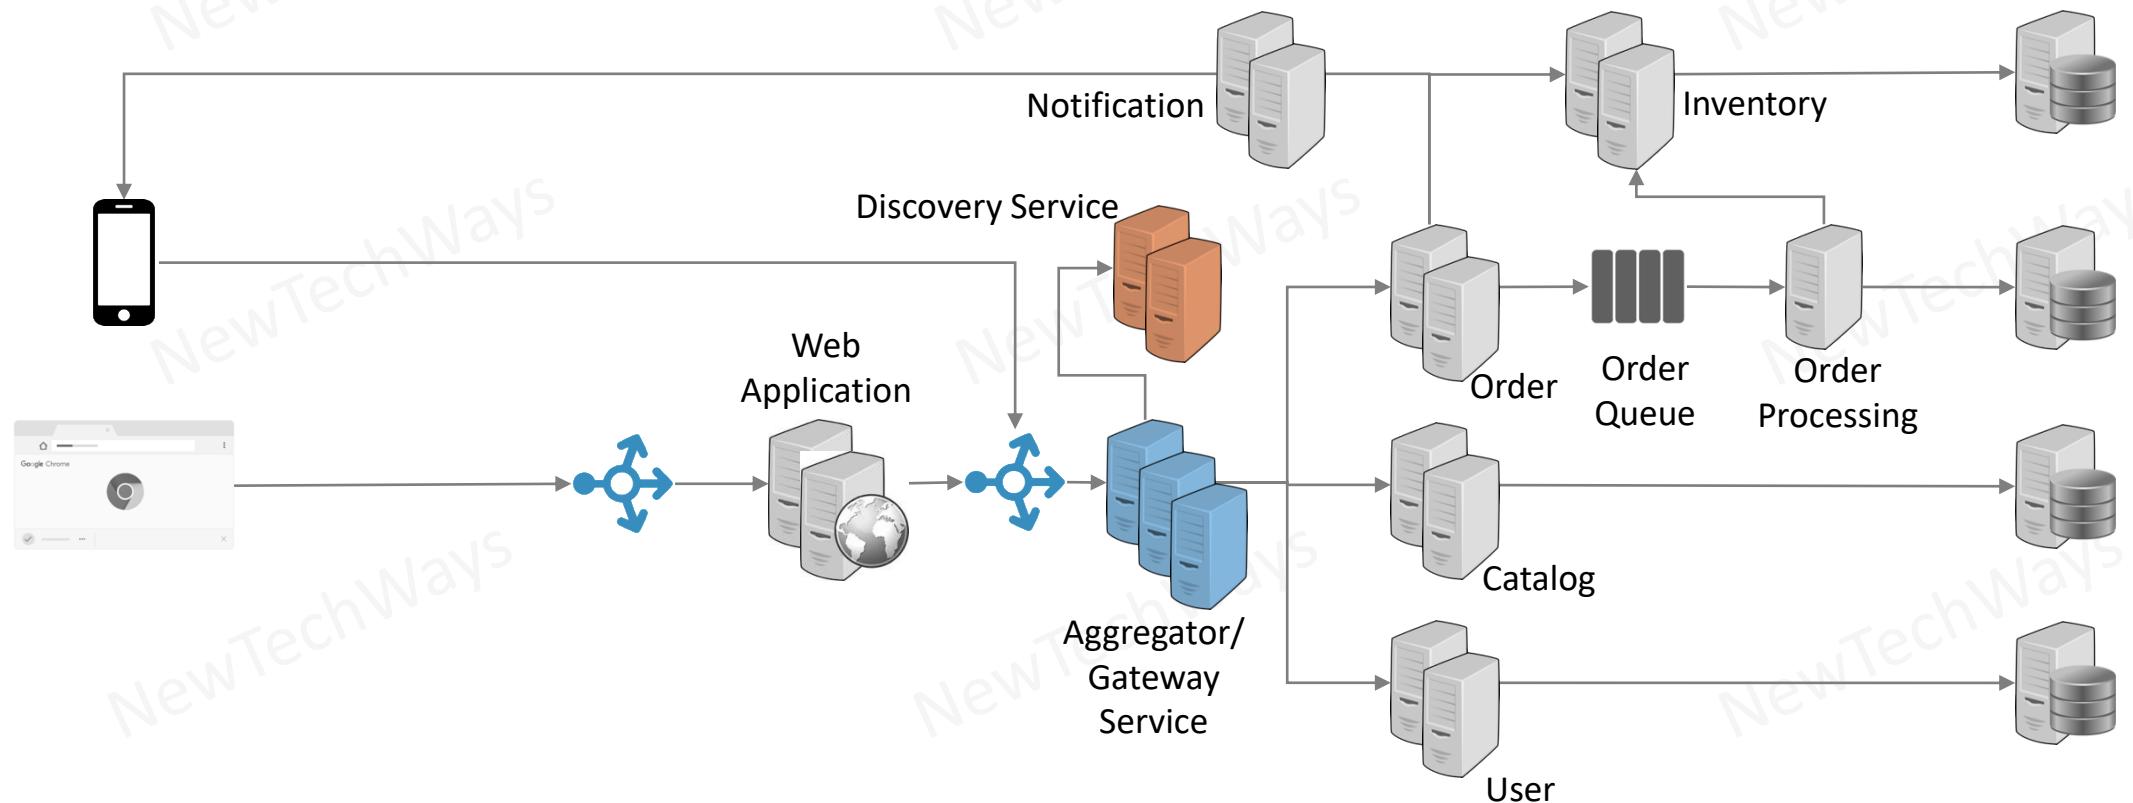
Stateless Replication

- Replication provides more computation power



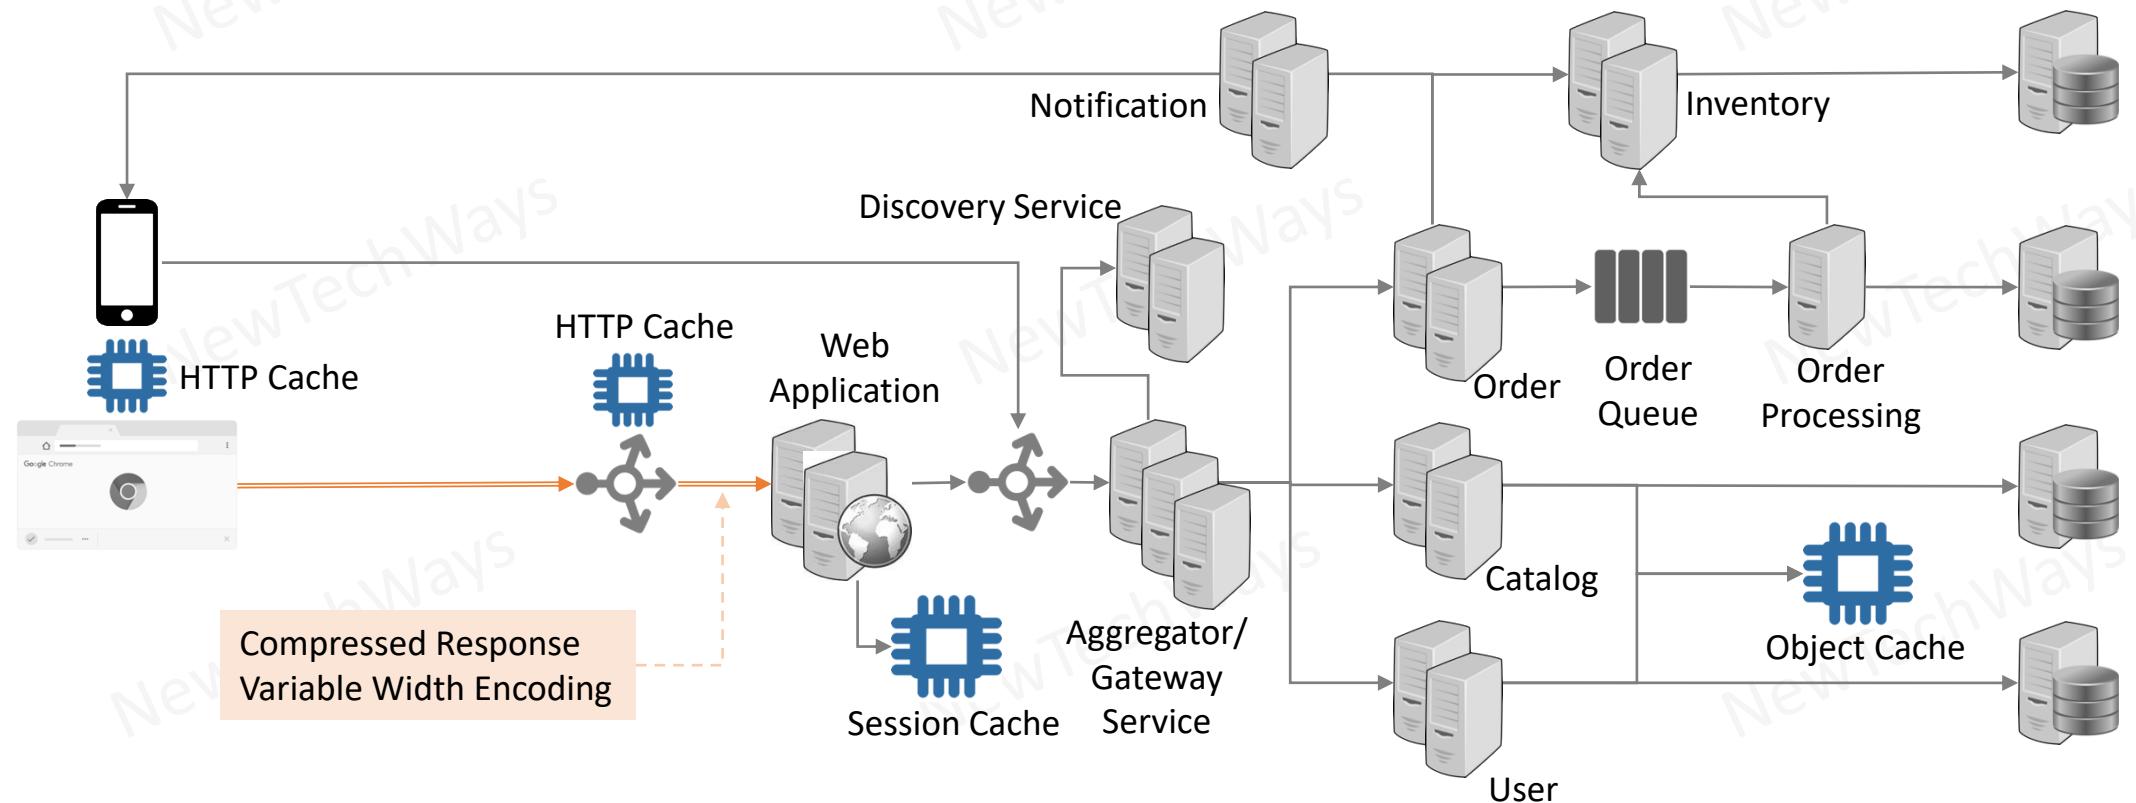
Load Balancing & Discovery

- Load Balancers for routing and load balancing
- Discovery services for naming and tracking healthy services



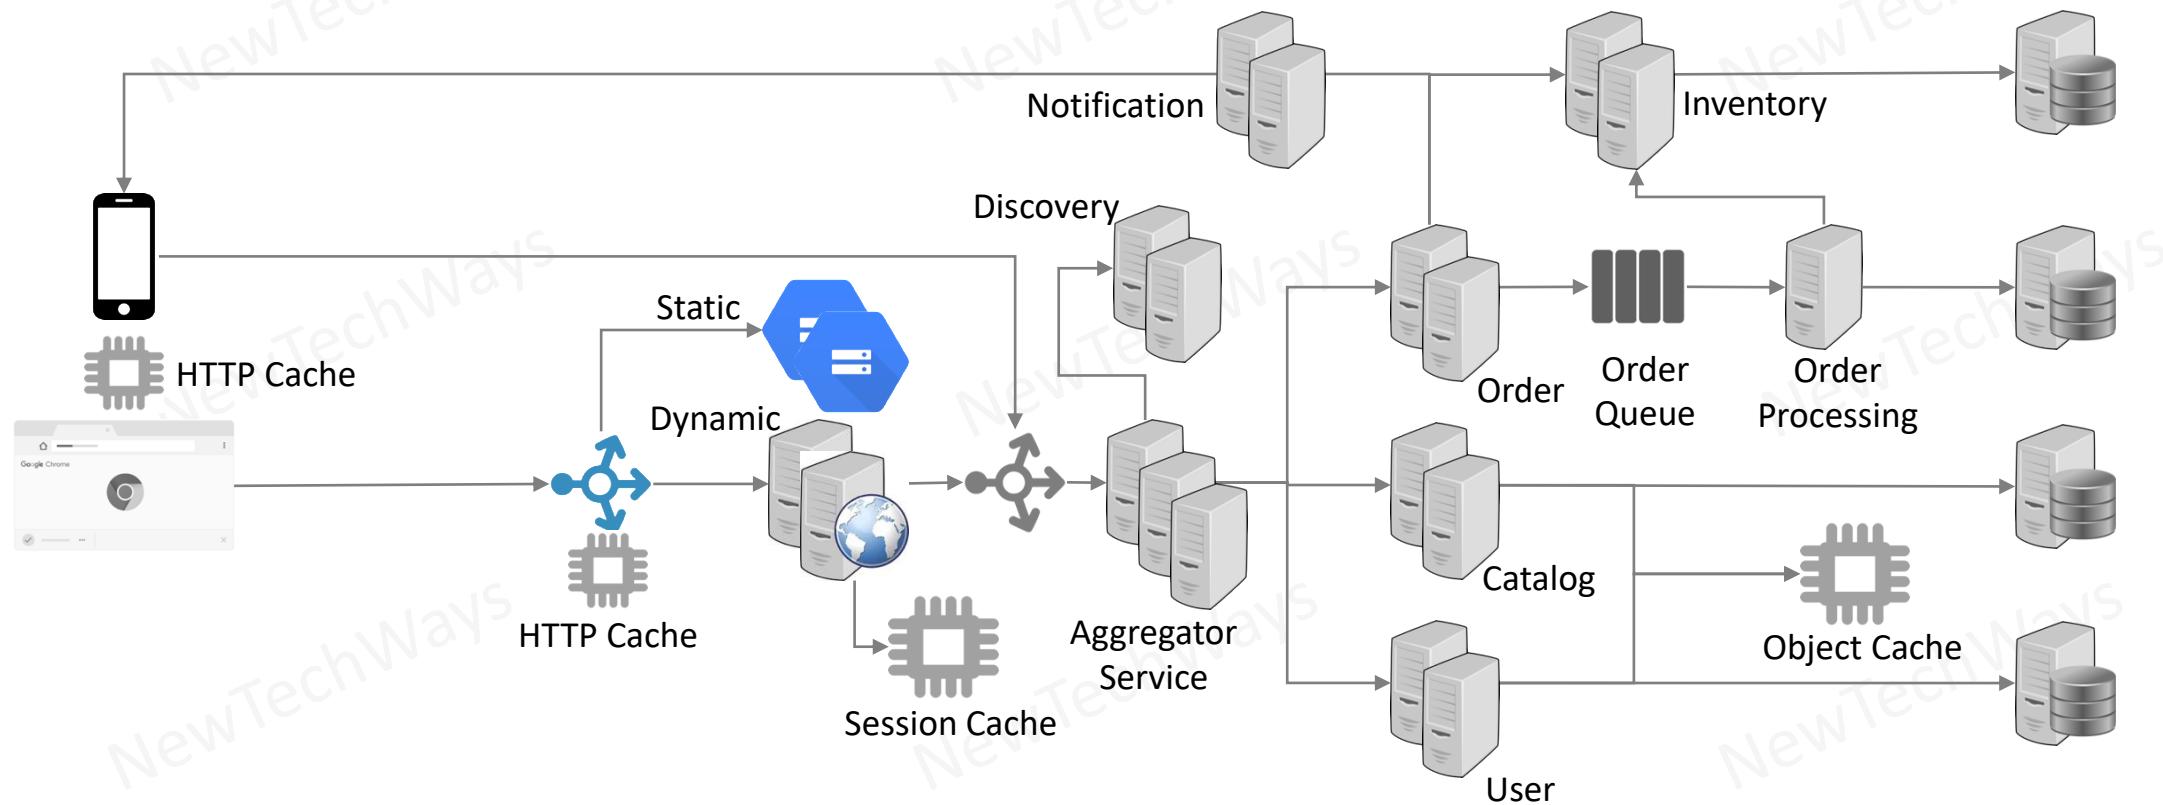
Caching & Compression

- Caching and compression for reducing response latency
- Caching can also reduce the backend load



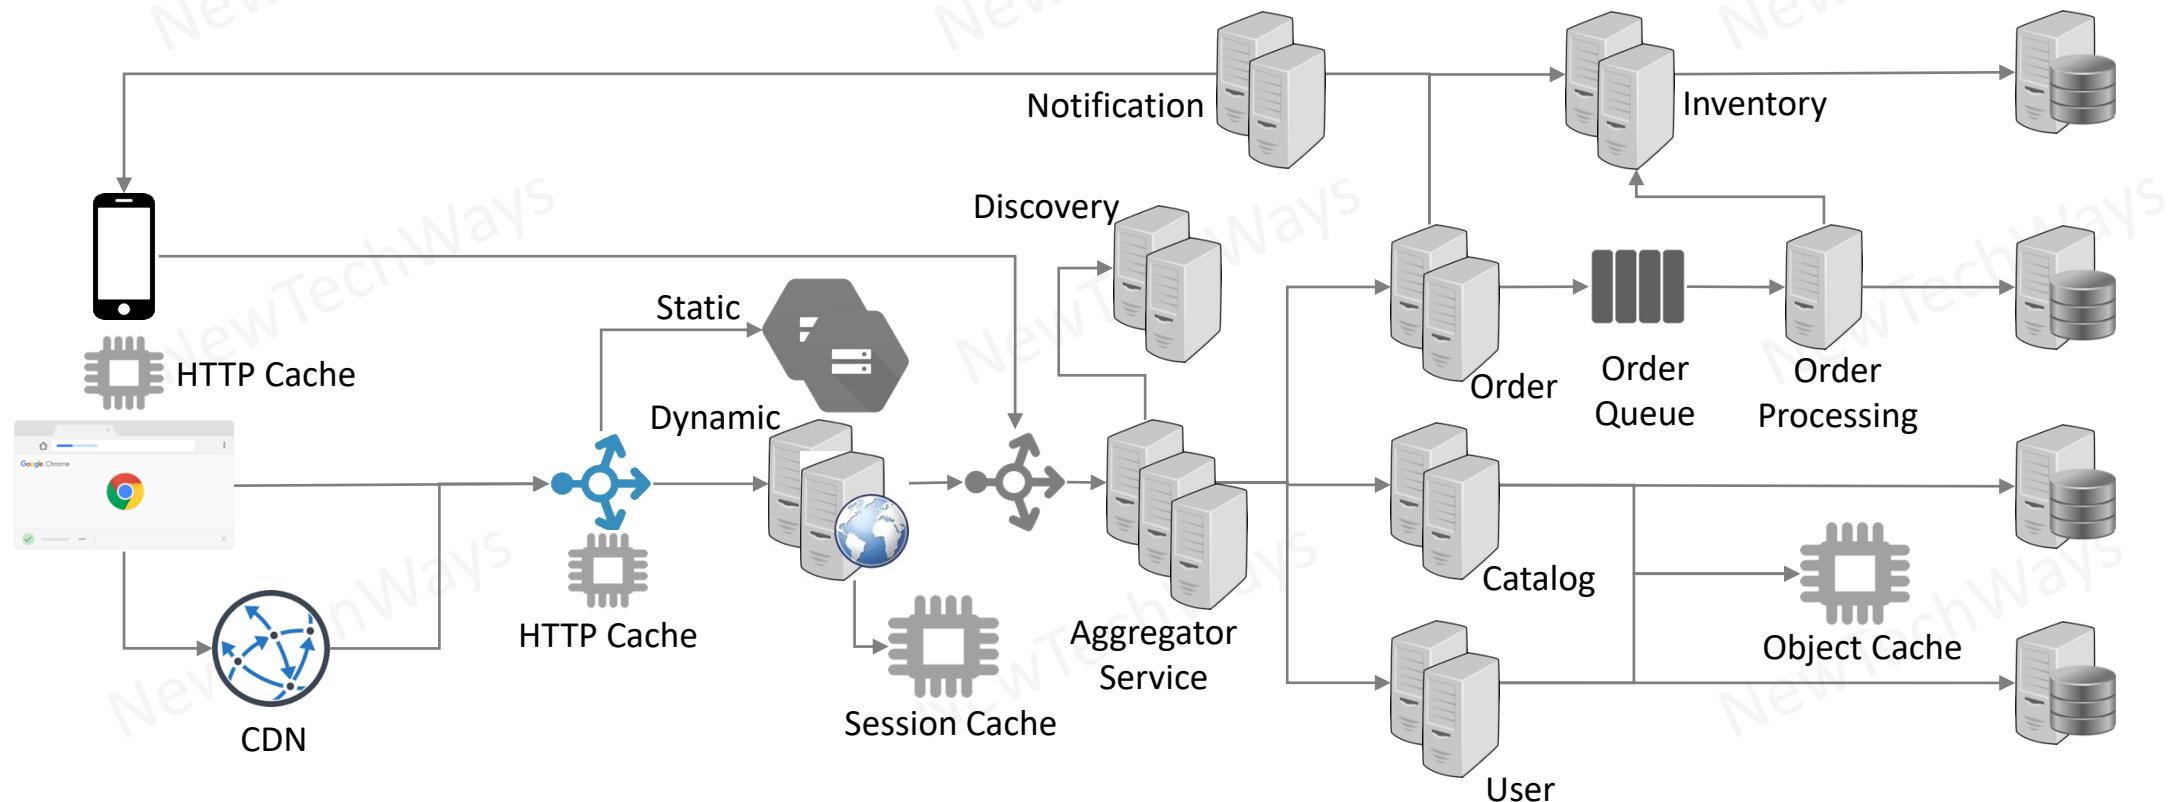
Serving Static Data

- High throughput Cloud Object Stores for serving static content



CDN & Edge Proxy

- CDN for global distribution of static content
- Edge Proxies for reducing latency through Early Termination, SSL Termination



Think About It!

27

- User Session Management
- Load Balancer
 - Routing & Sessions
 - Types – Software/Hardware, External/Internal
 - Policies
- Balancing Latency & Throughput
- RDBMS Scaling Limitations
- NoSQL Vs RDBMS
- Transaction Concurrency
- Eventual Consistency
- Routing for Global Scalability

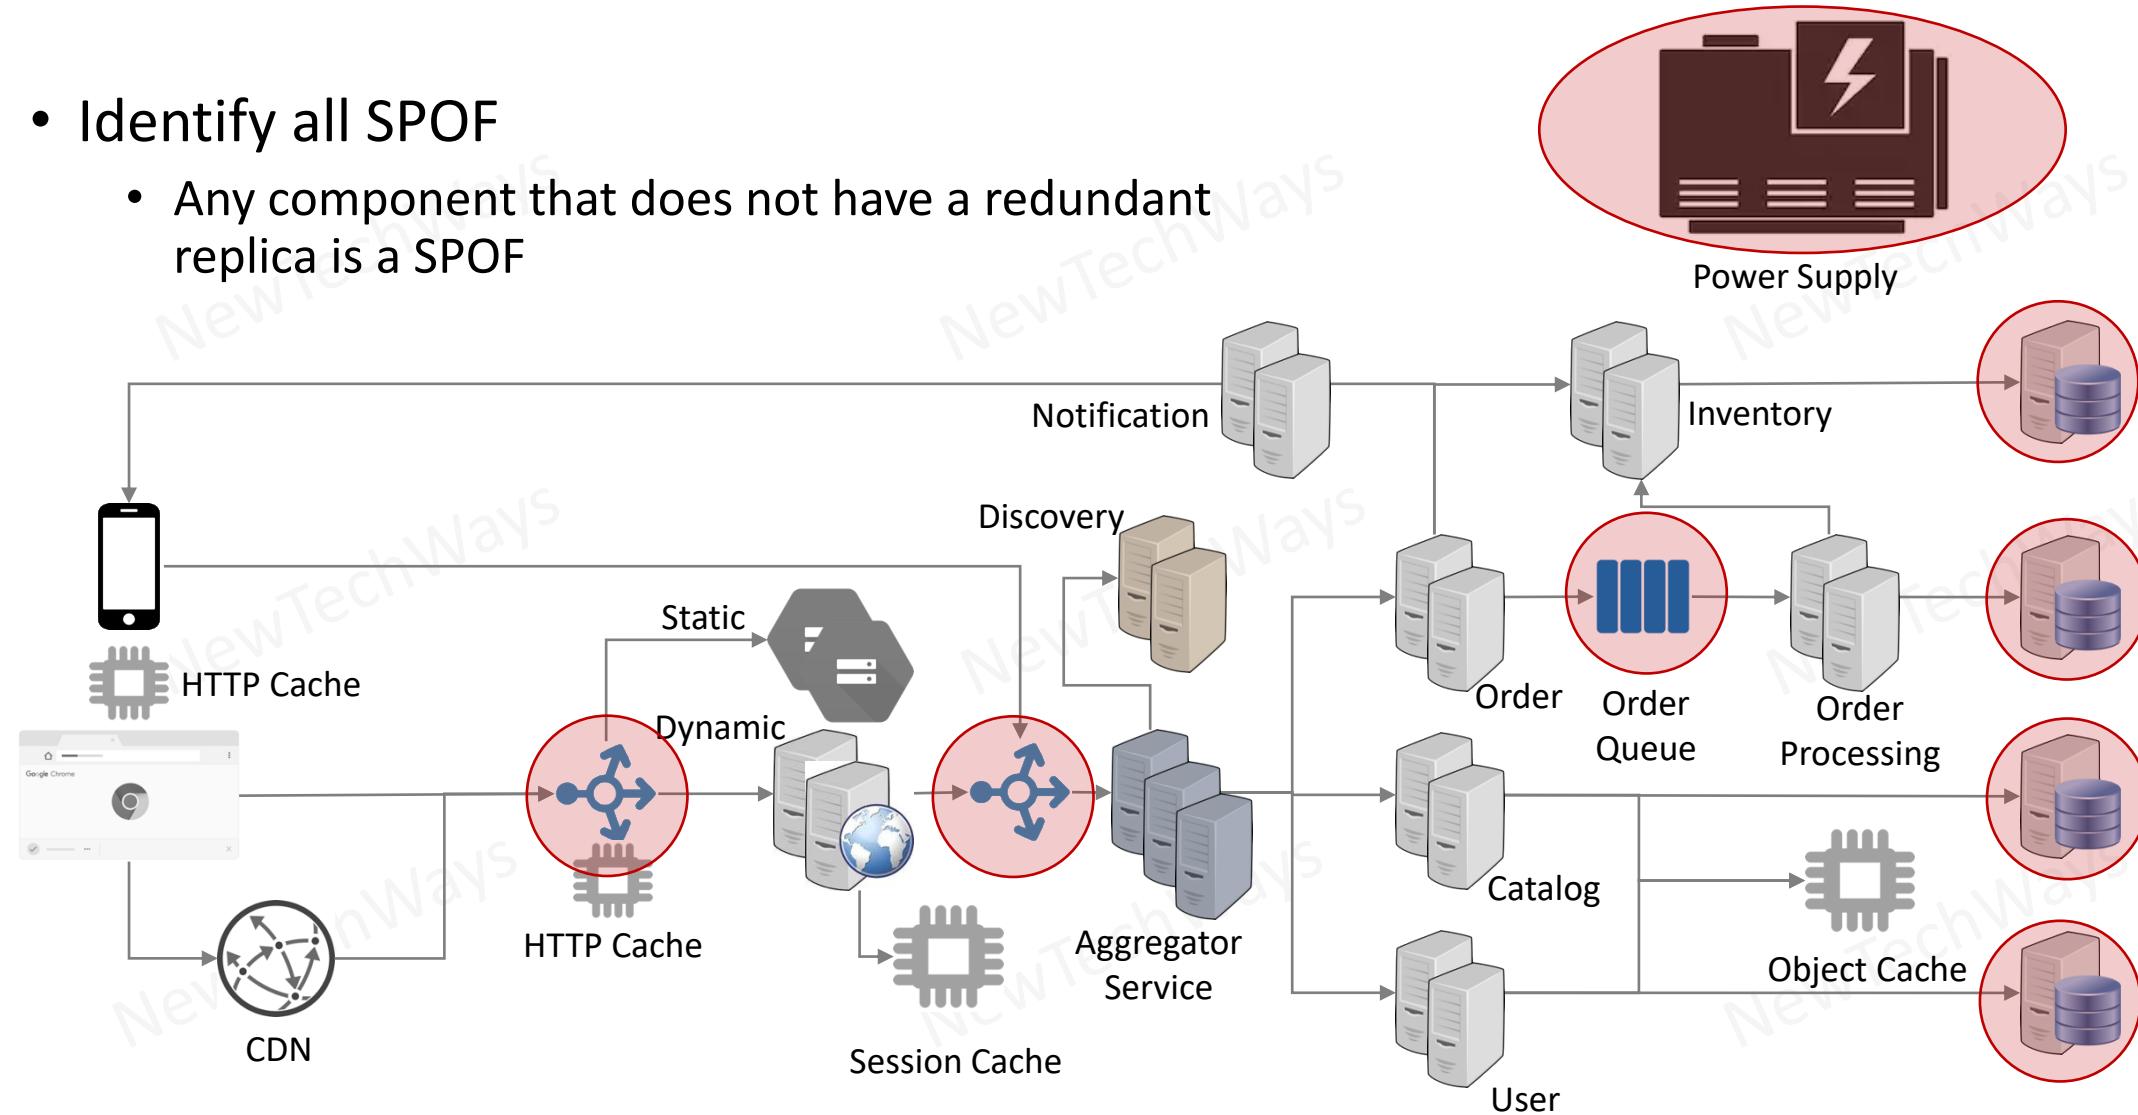
Reliability Principle

- Reliability
 - Normal functioning even in the presence of faults
- Availability
 - Always available even in the presence of faults
- Reliability and Availability are achieved mainly through Fault Tolerance
- Fault Tolerance requires
 - Provisioning Redundancy
 - Stateless & Stateful
 - Active, Passive, Cold
 - Fault Detection Mechanisms
 - Health-checks, heart-beats
 - Recovery or Failover
 - Stateless & Stateful



Single Point Of Failures

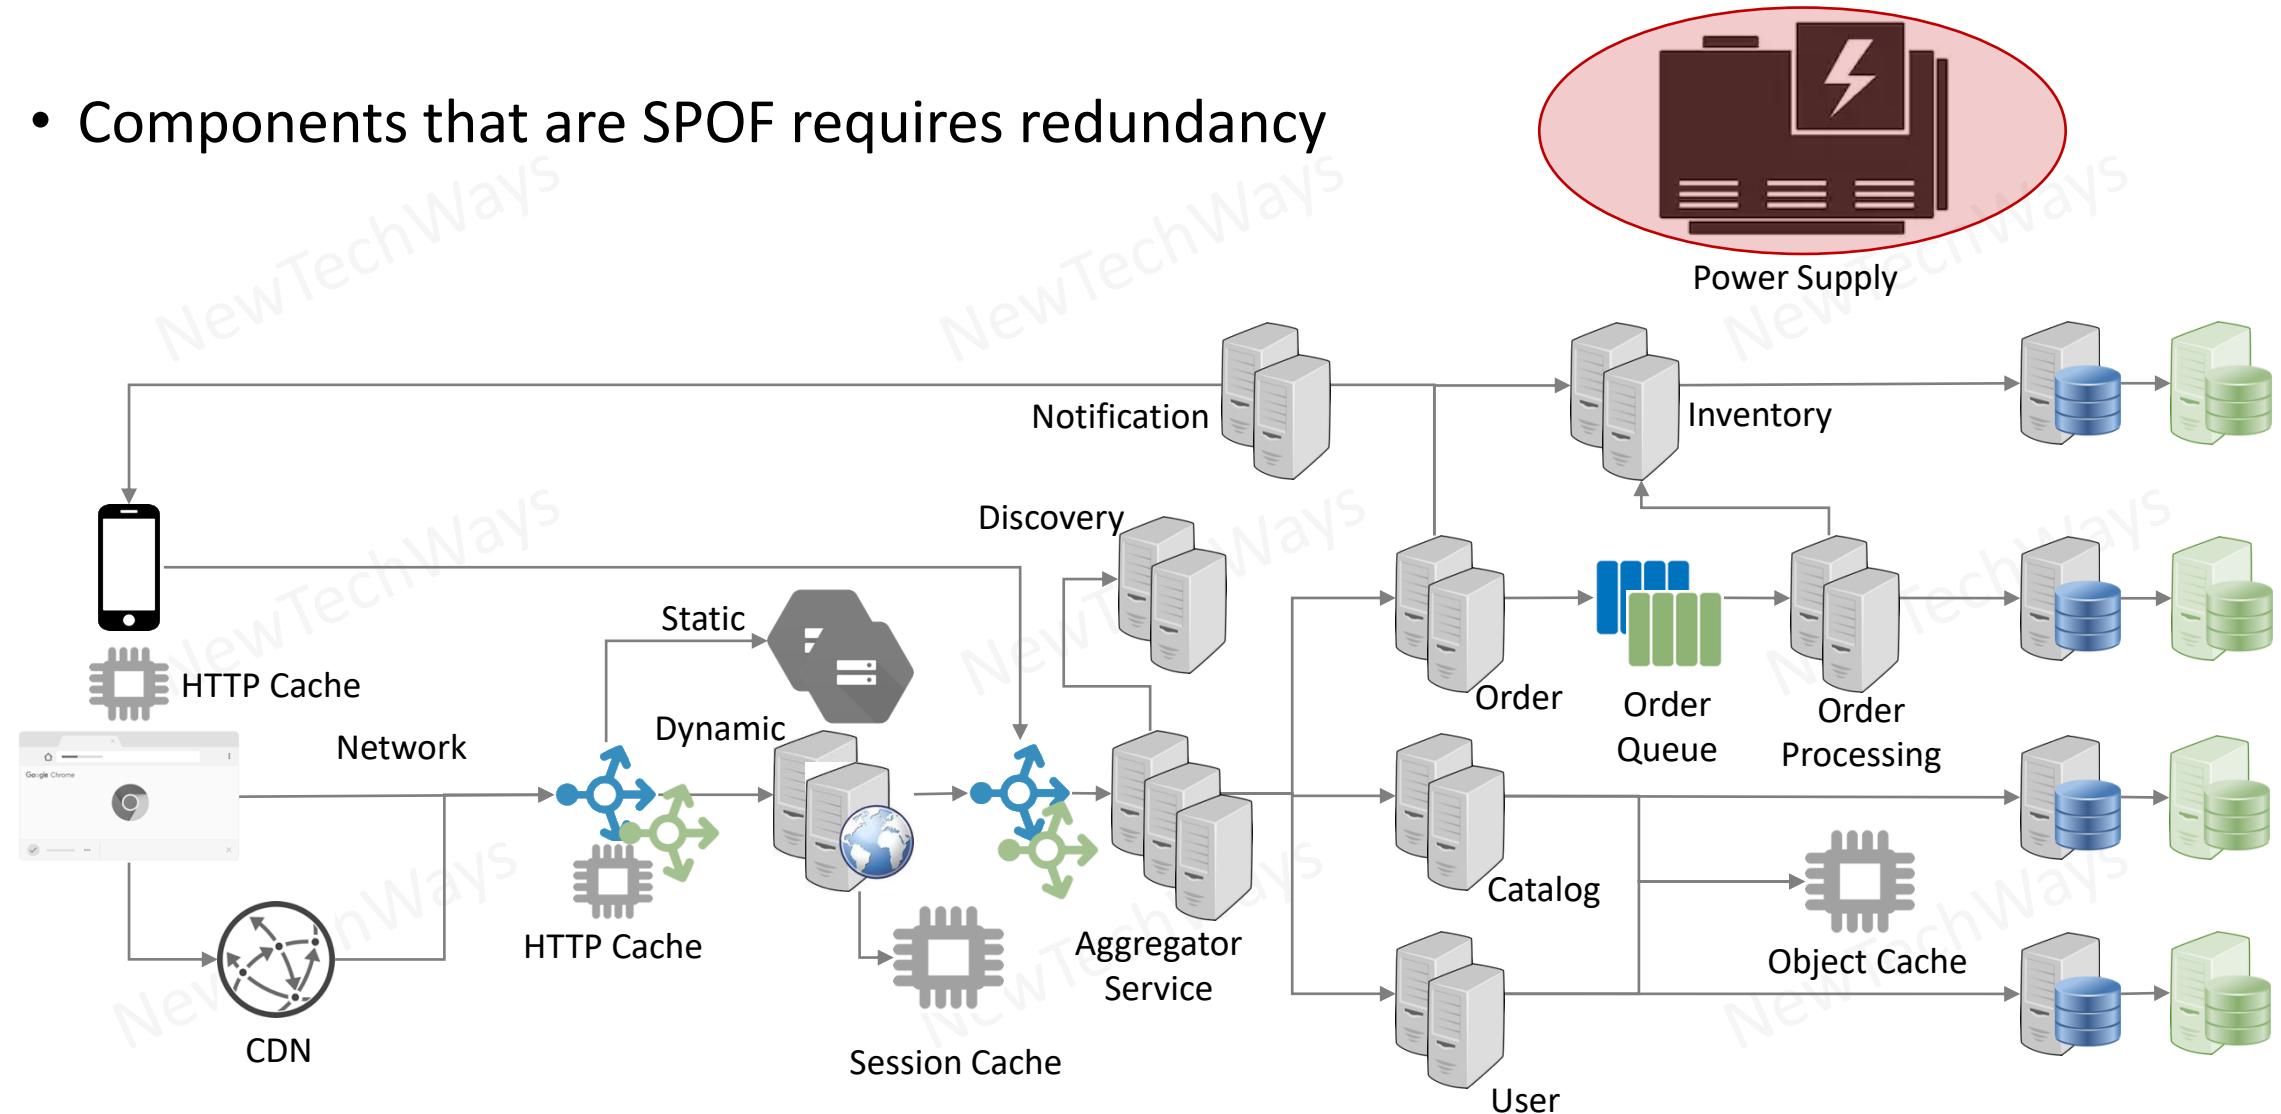
- Identify all SPOF
 - Any component that does not have a redundant replica is a SPOF



Redundancy

30

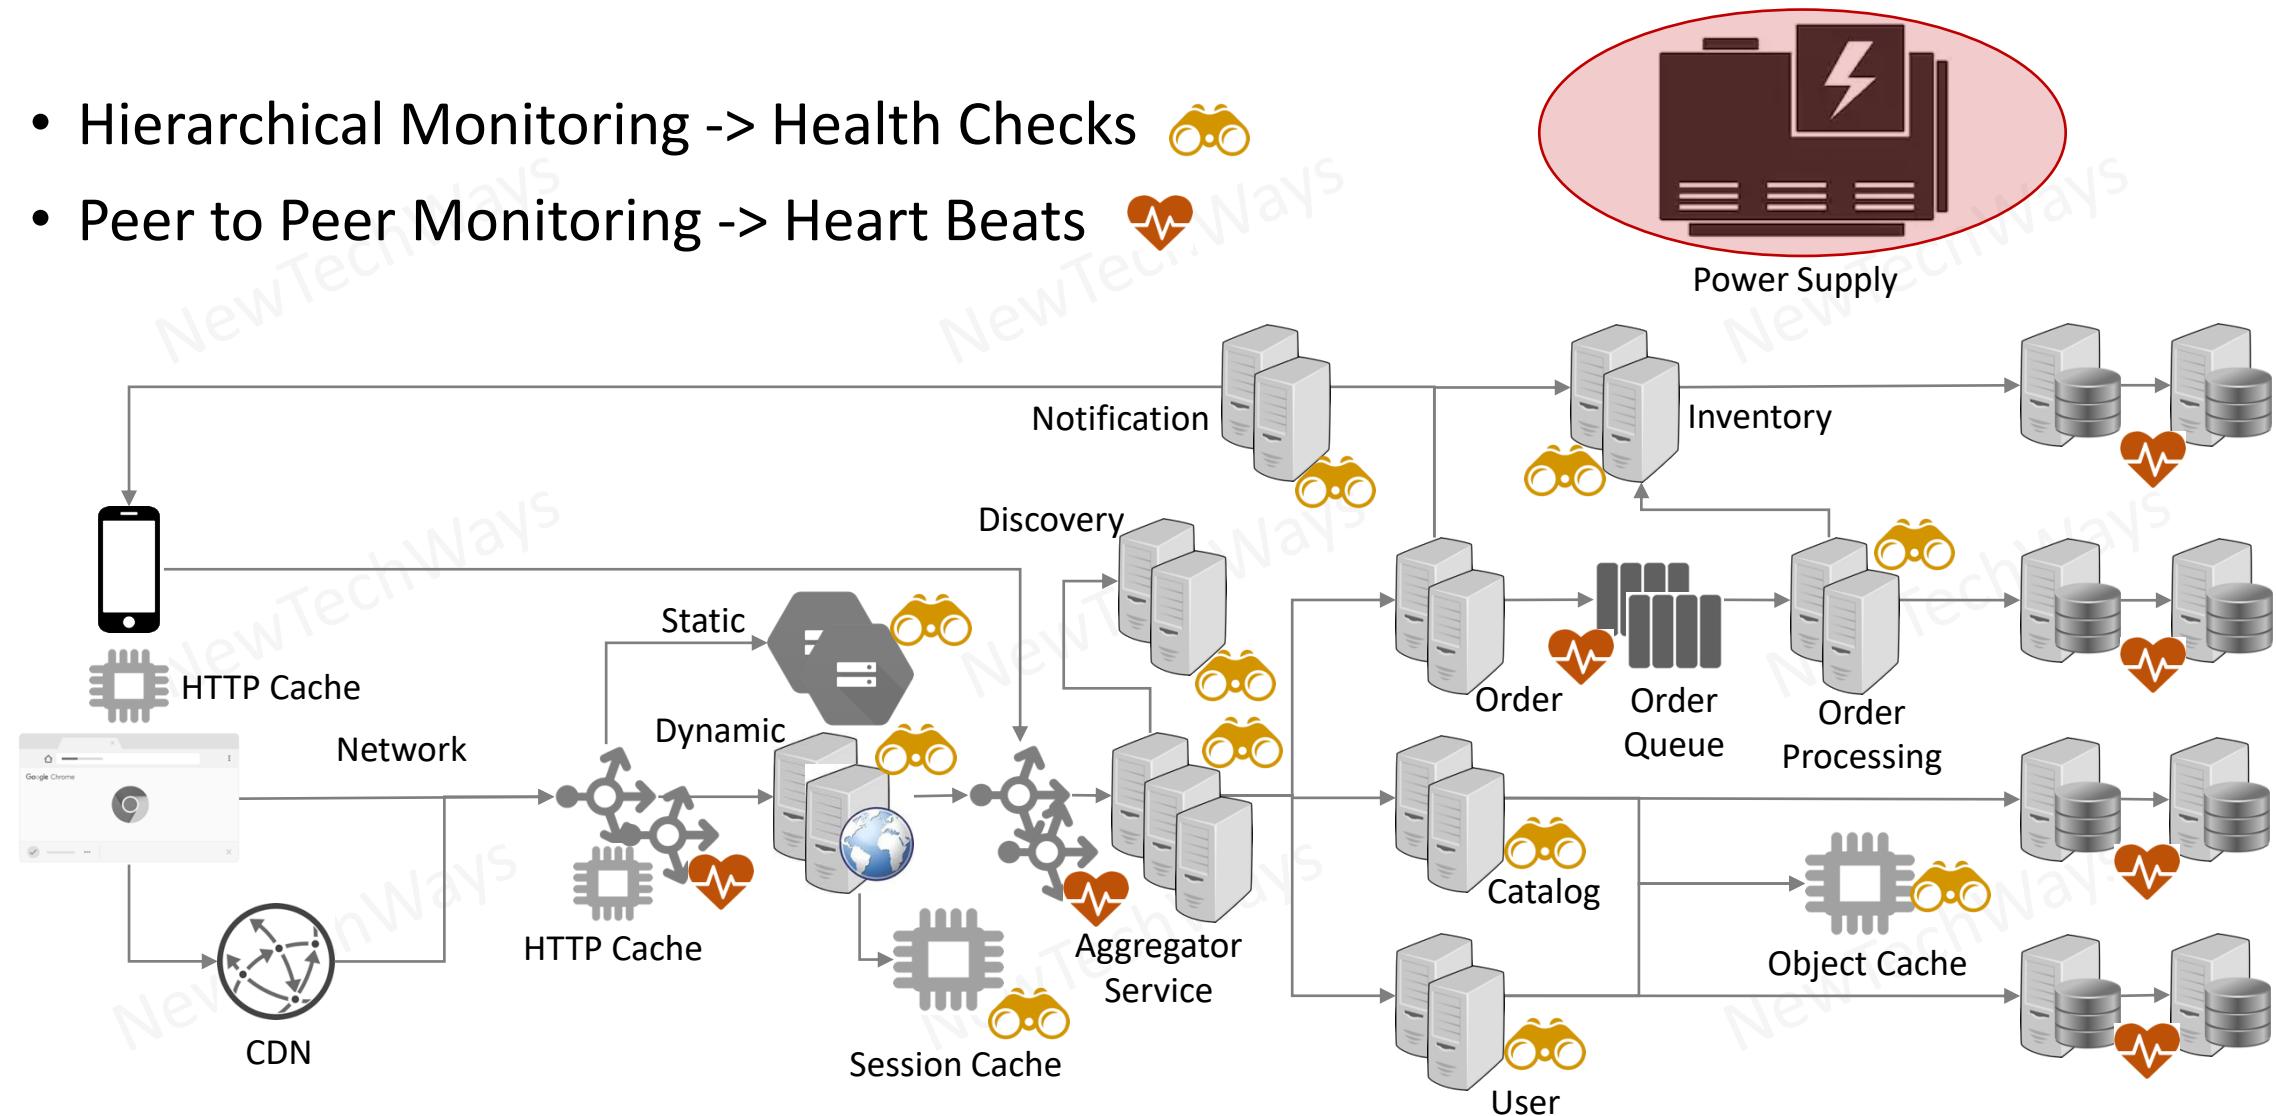
- Components that are SPOF requires redundancy



Monitoring For Fault Detection

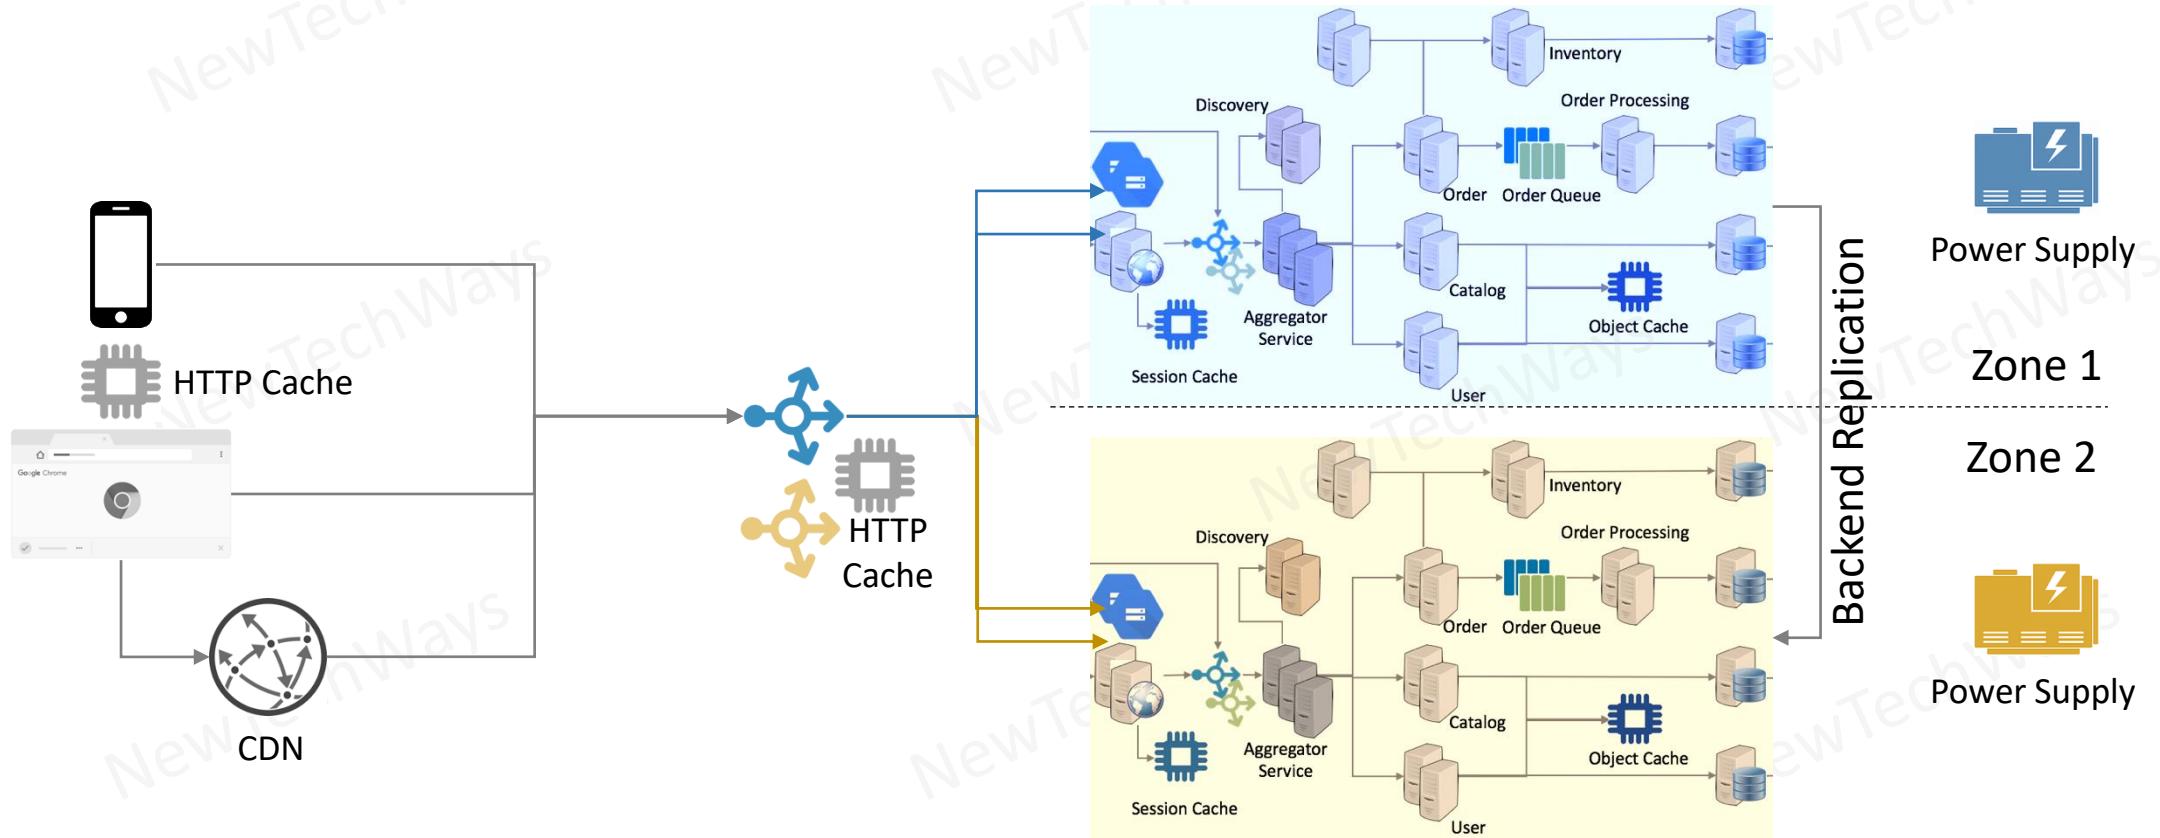
31

- Hierarchical Monitoring -> Health Checks 🕵️
- Peer to Peer Monitoring -> Heart Beats 💬



Zonal Redundancy

- Zonal redundancy for fault isolation



Think About It!

33

- Zonal Redundancy & Routing
- Regional Redundancy & Routing
- Redundancy & Recovery
 - Stateless Components
 - Stateful Components
- Automated Recovery
 - Detecting Faults
 - Triggering Failover
- Reliability Practices
 - Retries & Idempotency
 - Graceful handling of errors
 - Controlled service degradation
 - Load Shedding
 - Backpressure

Security Principles

- Least Privilege
 - Minimalistic user access rights just enough to perform a task
- Minimize attack surface area
 - Reduce entry points by authorizing every API request or Activity
 - Allow only minimum set of roles for each API
- Defense In Depth
 - Multiple layered resource access controls
- Separation of Duties
 - Different entities should have different roles
- Fail Secure
 - Failure should not expose more than required information through error messages or logs
- Weakest Link
 - Any system is as strong as its weakest link

OWASP Top 10 - 2017

A1:2017-Injection

A2:2017-Broken Authentication

A3:2017-Sensitive Data Exposure

A4:2017-XML External Entities (XXE) [NEW]

A5:2017-Broken Access Control [Merged]

A6:2017-Security Misconfiguration

A7:2017-Cross-Site Scripting (XSS)

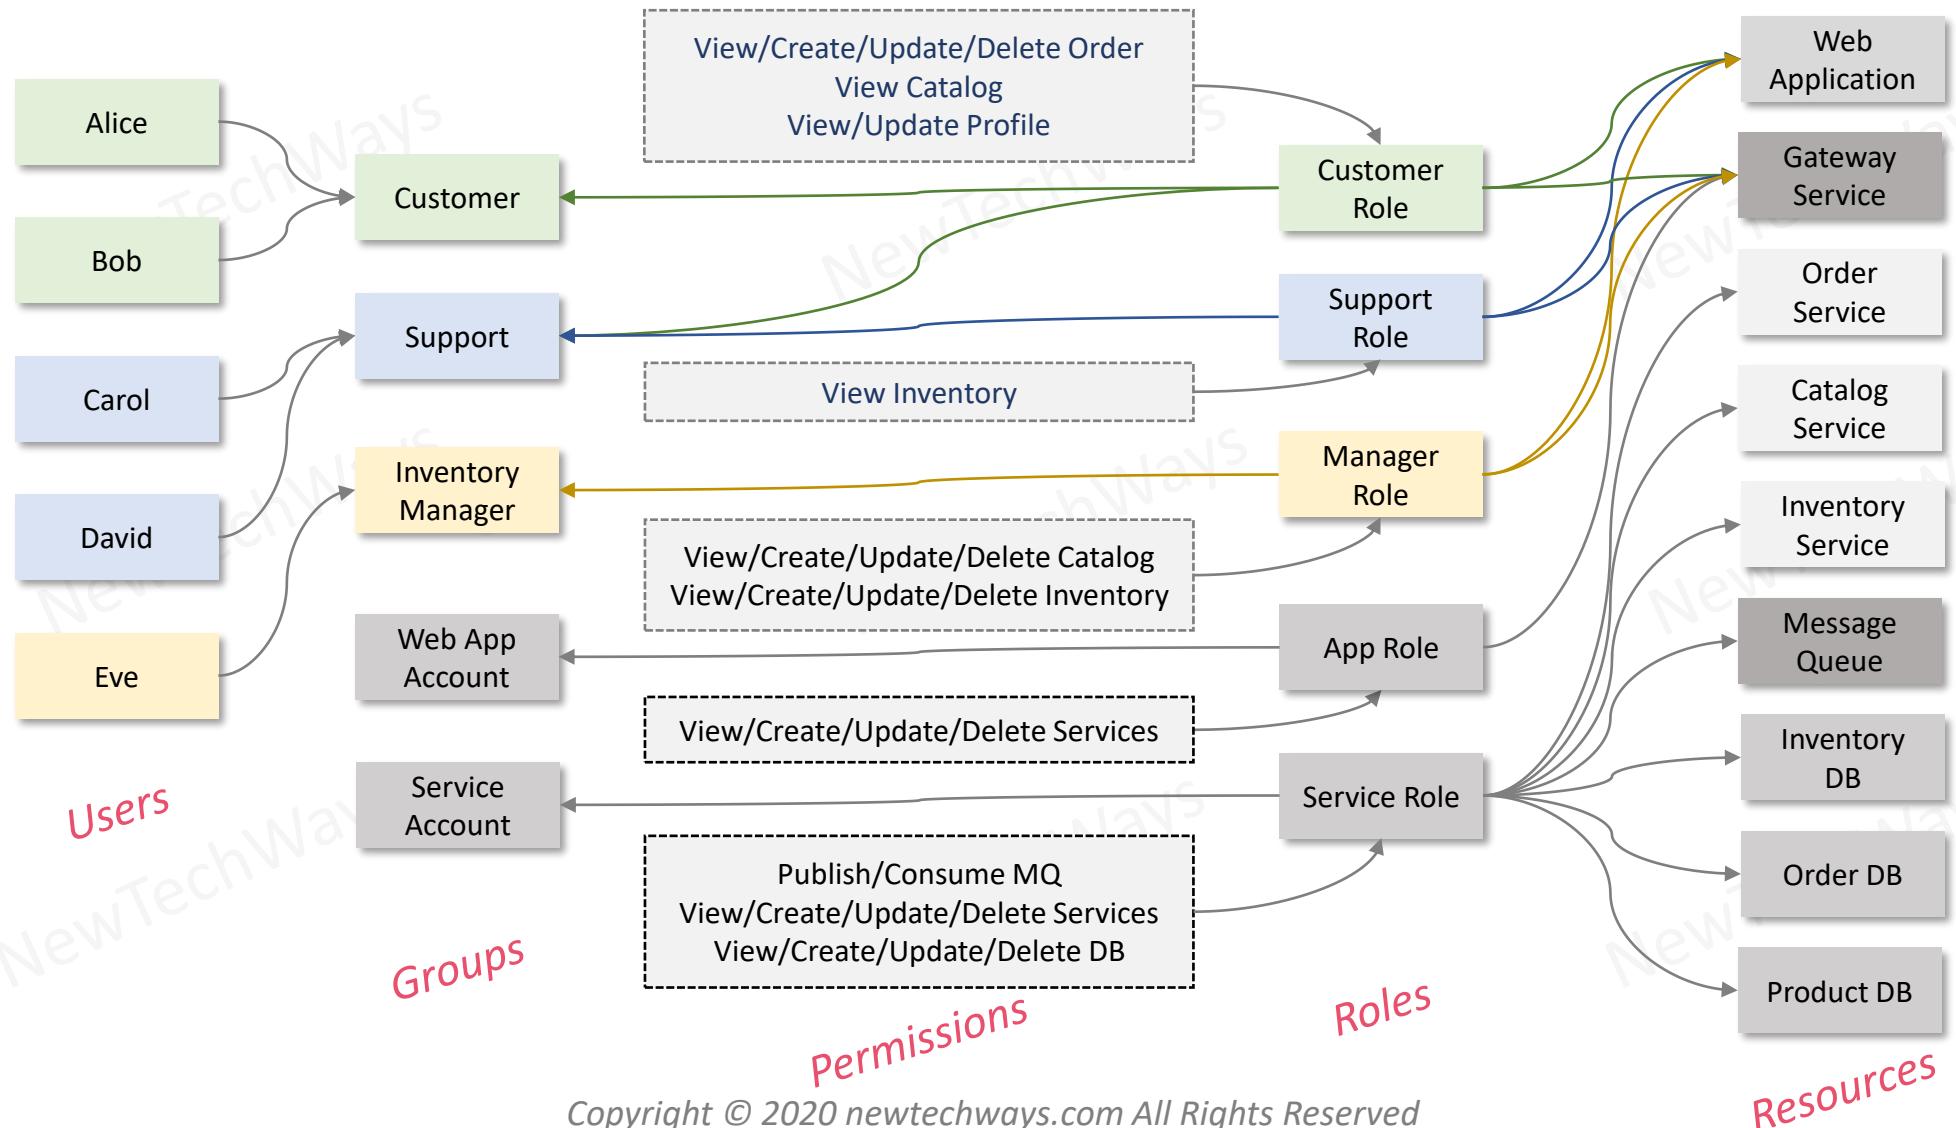
A8:2017-Insecure Deserialization [NEW, Community]

A9:2017-Using Components with Known Vulnerabilities

A10:2017-Insufficient Logging&Monitoring [NEW,Comm.]

Role Based Access Control

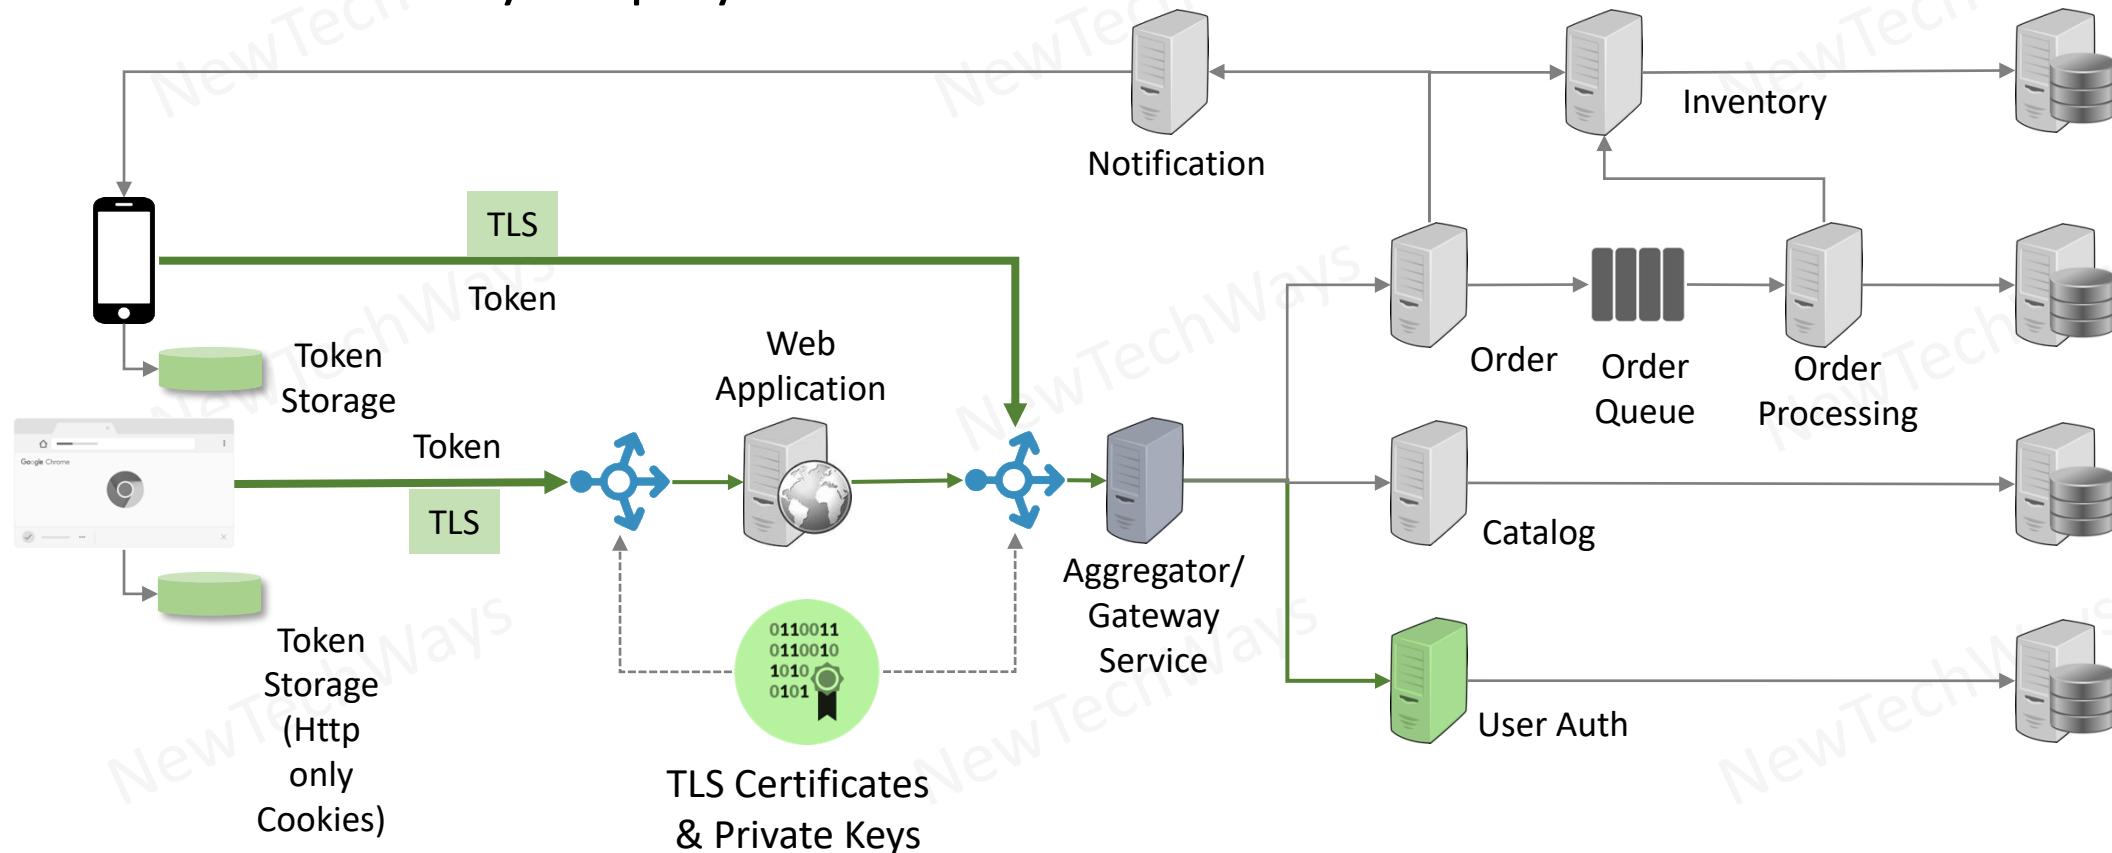
35



Authentication & Authorization

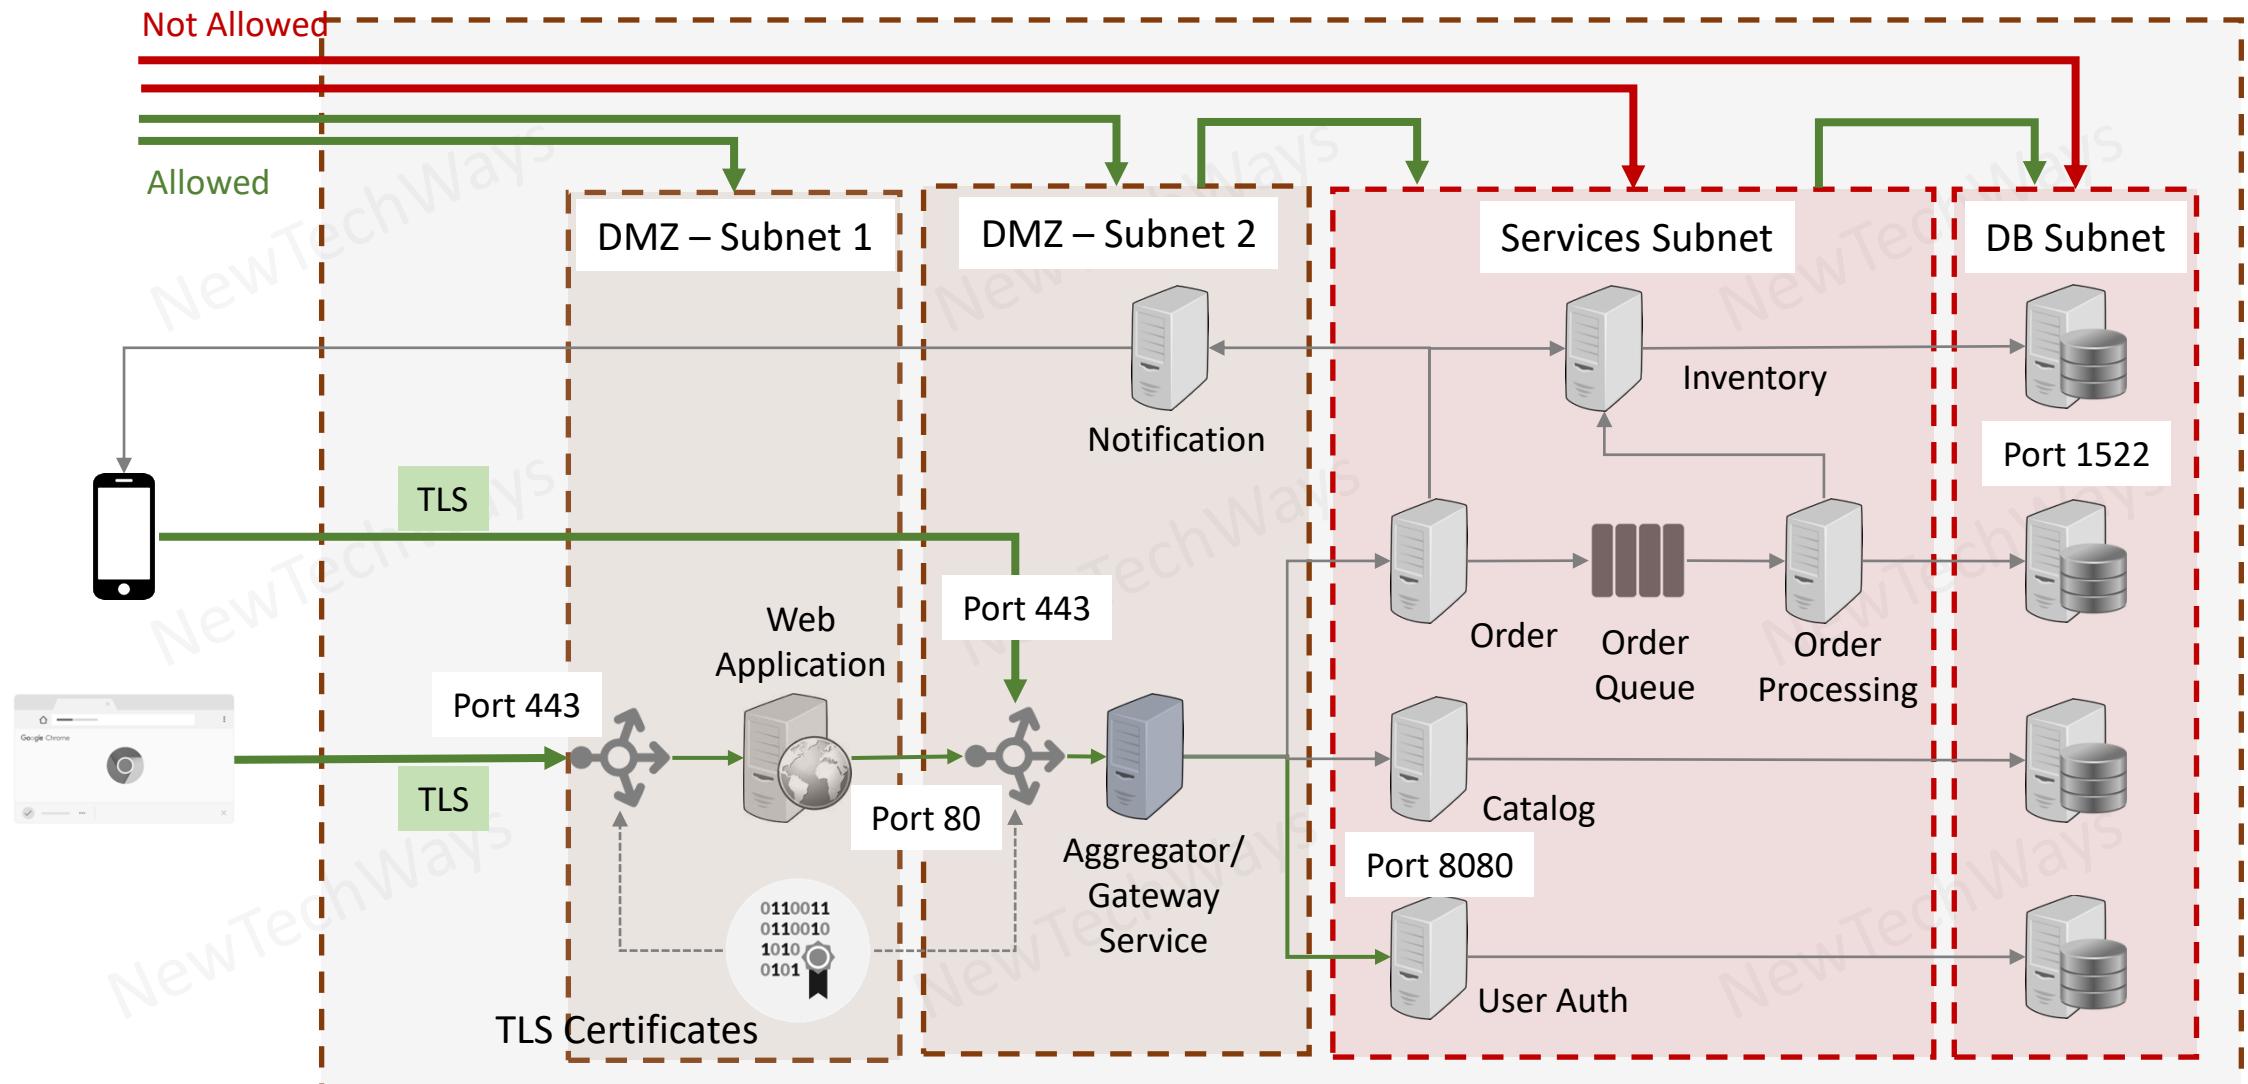
36

- Token based authentication over HTTPS
- Certificates & keys deployed on external load balancers



Network Security

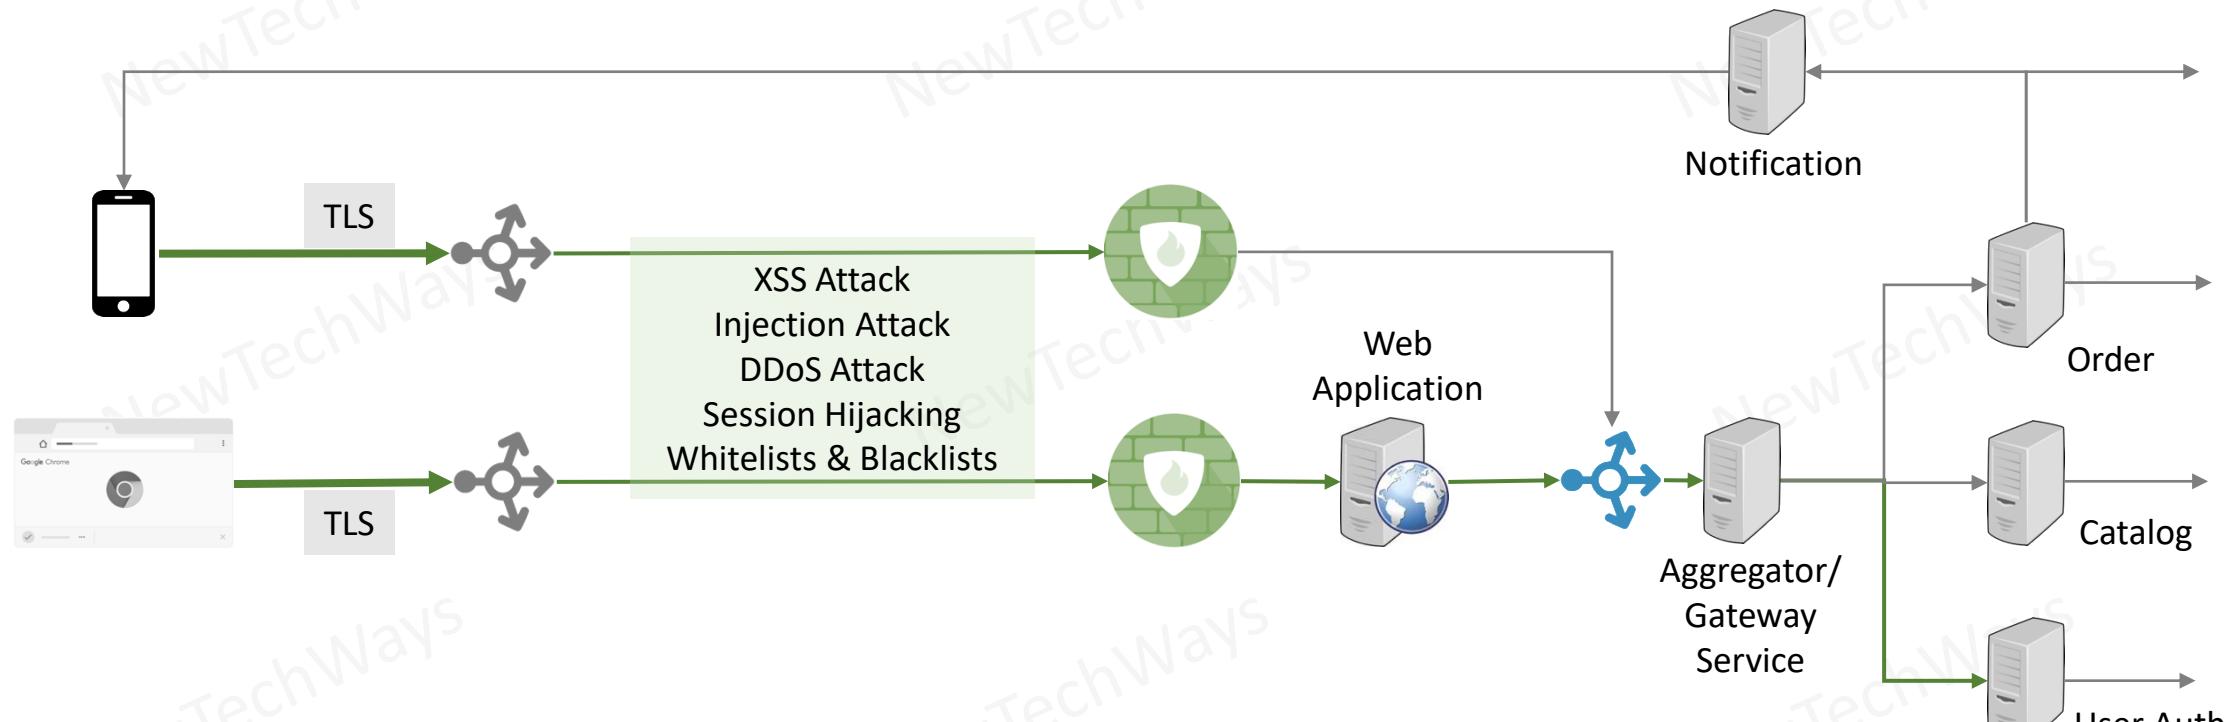
37



Web Application Firewalls

38

- Multiple security related attacks can be prevented by inspecting requests for common attacks and vulnerabilities by using Web Application Firewalls



Think About It!

39

- Security Token Storage
- Stateless & Stateful Authentication
- Third-Party Authorization with OAuth2
- Single Sign On
- TLS Termination
- Common Vulnerabilities
- Stored Data Encryption
- Key Management

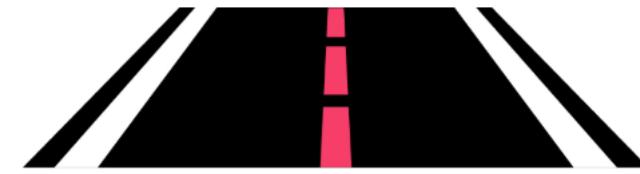
That's Just The Beginning

Further challenge lies in the detail
And we have sorted it out for you

If you are an Architect or a Developer
And you want to bridge the gap between
A Great Developer and a True Architect

You should get started with
'NewTechWays' course
'Developer To Architect'
And leapfrog your career

Thank You!



NewTechWays

*Academy of Software
Architecture*

Visit us at

<https://www.newtechways.com>